SEPARATING PREDICTED RANDOMNESS FROM RESIDUAL BEHAVIOR

Jose Apesteguia
ICREA-Universitat Pompeu Fabra
and Barcelona GSE

Miguel A. Ballester
University of Oxford

Abstract

We propose a novel measure of goodness of fit for stochastic choice models, that is, the maximal fraction of data that can be reconciled with the model. The procedure is to separate the data into two parts: one generated by the best specification of the model and another representing residual behavior. We claim that the three elements involved in a separation are instrumental in understanding the data. We show how to apply our approach to any stochastic choice model and then study the case of four well-known models, each capturing a different notion of randomness. We illustrate our results with an experimental data set. (JEL: D00)

> **Teaching Slides**
> A set of Teaching Slides to accompany this article are available online as Supplementary Data.

1. Introduction

Choice data arising from either individual or population behavior are often probabilistic in character. Currently, there is a renewed interest in finding better methods for modeling stochastic behavior, and the literature offers a battery of models incorporating randomness in various different ways.[1] This paper discusses a novel goodness-of-fit

---

1.   Recently published papers include Gul and Pesendorfer (2006), Dickhaut, Rustichini, and Smith (2009), Caplin, Dean, and Martin (2011), Ahn and Sarver (2013), Gul, Natenzon and Pesendorfer (2014), Manzini and Mariotti (2014), Fudenberg and Strzalecki (2015), Fudenberg, Iijima and Strzalecki (2015), Matejka and McKay (2015), Barseghyan, Molinari and O'Donoghue (2016), Brady and Rehbeck (2016),

measure for stochastic choice models, namely, the (tight) upper bound in the portion of the data that can be reconciled with the model. This approach requires us to separate the data into two parts: one that is generated by a particular specification of the model and is to be maximized, and the remaining unstructured portion, which is to be minimized. We refer to the first part as predicted randomness and to the second as residual behavior.

The separation exercise highlights three key elements. The first is the maximal fraction of data explained by the model, which we take as our goodness-of-fit measure. This indicates the ability of the model to explain actual behavior. The second is an optimal specification of the model, which, if it explains a large portion of the stochastic data, is obviously a potentially useful tool in counterfactual scenarios, such as those associated with prediction problems. The third is a description of residual behavior, which aids understanding of the relationship between actual behavior and the choice model, because it endogenously enables the identification of the menus and choices for which the model deviates farthest from the data. This information may be relevant when it comes to revising a model.

More formally, given the grand set of alternatives X, SCF denotes the set of all stochastic choice functions, that is, all possible descriptions of the choice probabilities of each alternative in each menu within the domain. The aim is to explain data ρ, which is taken to be a stochastic choice function, in light of model Δ, which is defined as a collection of stochastic choice functions. Model Δ describes all the possible predictions the analyst considers relevant. For example, it may encompass all the parametric specifications of the analyst's preferred choice model, including those accounting for measurement error or unobserved heterogeneity. A triple $\langle \lambda, \delta, \varepsilon \rangle$, where $\lambda \in [0, 1]$, $\delta \in \Delta$, and $\varepsilon \in$ SCF, such that $\rho = \lambda\delta + (1 - \lambda)\varepsilon$ describes a possible separation of data ρ into a portion λ explained by the instance of the model δ and another portion $1 - \lambda$, which is the unstructured residual behavior ε. A separation is maximal if it provides the maximal value of λ. As well as showing that, for any closed model Δ, maximal separations always exist, Proposition 1 in Section 3 also characterizes their structure. The result shows that maximal separations are identified by the following maxmin operator, which begins by computing the minimum data-to-prediction ratio, across all observations, for every instance of the model. The solution is then given by the model instance that maximizes this ratio. This is a simple method, applicable to any model, and potentially instrumental in the analysis of particular models, as will be shown later.

Section 4 analyzes four well-known stochastic choice models, each predicting randomness in a very different way. In all four cases, we build on Proposition 1 in order to provide tailored results describing the structure of the maximal separations of the different models. By elaborating on the structural properties of the respective stochastic choice model, this exercise complements the conceptual understanding of the maximal separation approach, while also facilitating its practical implementation. We start with the paradigmatic decision-making model in economics: the deterministic choice model, where the individual always selects the alternative that maximizes a

---

Caplin and Dean (2015), Agranov and Ortoleva (2017), Apesteguia, Ballester, and Lu (2017), Apesteguia and Ballester (2018), Webb (2019), and Natenzon (2019).

preference relation, and hence there is no predicted randomness whatsoever. Thus, when a stochastic choice function is analyzed from the perspective of the deterministic model, any stochasticity in the data must be regarded as residual behavior. Given the overwhelming use of this model, it seems appropriate that it should be the first in our analyses of particular cases. Proposition 2 provides a simple recursive argument over the cardinality of the menus for computing the maximal separation of the deterministic model.

Next, we take three stochastic choice models, starting with the tremble model, where randomness represents the possibility of choice errors. In the tremble model, the decision maker maximizes a preference relation with probability $(1 - \gamma)$ and, with probability γ, randomizes over all the available alternatives. Proposition 3 describes how the technique developed for the deterministic model can be extended to this case. We then analyze the model proposed by Luce (1959), also known as the logistic model. The Luce model incorporates randomness in the utility evaluation of the alternatives. Proposition 4 gives simplicity to the analysis of the Luce model by showing that the observations yielding the minimum data-to-prediction ratio in a maximal separation obey a particular structure. Finally, we study a class of random utility models (RUMs) incorporating randomness in the determination of the ordinal preference that governs choice. In particular, we study the class of single-crossing RUMs (SCRUMs; Apesteguia, Ballester, and Lu 2017), which has the advantage of providing tractability, while also being applicable to a variety of economic settings. Proposition 5 gives the corresponding maximal separation, by applying a recursive argument over the collections of preferences, which are in the support of the RUM.

Section 5 reports on an empirical application of our approach. We use a preexisting experimental data set comprising 87 individuals faced with binary menus of lotteries. We take the aggregate data for the entire population and illustrate the practicality of our theoretical results, obtaining the maximal separation results for all the models discussed in this paper. We first show that the maximal fraction of the data explained by the deterministic model, that is, its goodness of fit, is 0.51, and that the preference relation identified in the maximal separation basically ranks the lotteries from least to most risky. The tremble model identifies exactly the same preference relation, together with a tremble probability of 0.51, which increases the fraction of data explained to 0.68. The Luce model also increases the fraction of data explained to 0.74, and identifies a utility function over lotteries that is ordinally close to the preference ranking of the deterministic and tremble models. Finally, we implement the SCRUM, assuming the utility functions given by constant relative risk aversion (CRRA) expected utility. We obtain that the fraction of data explained increases further to 0.78, with the largest mass being assigned to a preference exhibiting high levels of risk aversion.

Section 6 compares the maximal separation approach with other goodness-of-fit measures, such as maximum likelihood and least squares. We argue that, by focusing on the largest deviations from the data, maximal separation is particularly accurate in predicting low probabilties. We then use the experimental data set to illustrate this point empirically, and thus confirm the existence of important complementarities between the maximal separation technique and standard techniques, when seeking a deeper understanding of the data.

Section 7 concludes by discussing three aspects of the maximal separation approach. We first briefly analyze the model-selection issue by discussing the case of an analyst wishing to compare the maximal fractions of data explained by different models. Secondly, we comment on the pros and cons of imposing further technical structure on the stochastic choice models. Finally, we consider the case in which the notion of maximal separation is slightly modified by restricting the space of possible residual behavior, which can be combined with the predicted randomness given by a model, and conclude by suggesting some potentially fruitful ways of interpreting residual behavior.

2. Related Literature

Rudas, Clogg, and Lindsay (1994) developed a novel proposal in statistics, in what is now known as the mixture index of fit for contingency tables. Given a multivariate frequency distribution, Rudas, Clogg, and Lindsay (1994) suggest measuring the goodness of fit of a given model using a two-point mixture, which entails calculating the largest fraction of the population for which a distribution belonging to the model fits the data, while leaving the complementary fraction as an unstructured distribution. Rudas (1999) extends the use of the mixture index to continuous probability distributions, and relates the optimal solution to minimax estimation.[2] The maximal separation technique imports the same logic for the study of stochastic choice functions, which differ from contingency tables in that they involve collections of interrelated probability distributions, one for each available menu of alternatives, where the interrelation is choice model-dependent. Interestingly, Böckenholt (2006) claims that new methodologies are needed to understand the systematic behavioral violations of RUMs, and, without elaborating, suggests the mixture index of fit as a potential tool for this purpose. In this paper, we undertake this challenge by extending the methodology, not only to RUMs, but to every possible stochastic choice model, and then incorporate these ideas into decision theory and economics.

In economics, Afriat (1973) made the first in a long history of proposals for indices that measure the consistency of revealed preferences with the deterministic, rational model of choice. Afriat's suggestion for a consumer setting was to compute the minimal monetary adjustment required to reconcile all observed choices with the maximization of some preference, an idea later generalized by Varian (1990). Alternative suggestions by Houtman and Maks (1985), and more recently by Dean and Martin (2016), are to compute the maximal number of data points that are consistent with the maximization of some preference. Apesteguia and Ballester (2015) and Halevy, Persitz, and Zrill (2018) suggest consistency measures to compute the minimal welfare loss from inconsistent choices with respect to some preference. Relevantly, Apesteguia and Ballester (2015) show axiomatically that all these measures have a

---

2. The statistical literature offers a number of applications of these ideas, and develops algorithms for the implementation of the mixture index to contingency tables (see, e.g., Dayton, 2003; Liu and Lindsay, 2009).

common structure, and search for a preference that minimizes a given loss function, ultimately providing both a goodness-of-fit measure and the best possible description of behavior.[3] The maximal separation approach shares the spirit of all these consistency measures, because it also provides a goodness-of-fit measure and the best description of behavior when applied to the deterministic rational model. In Appendix D, we formally compare the maximal separation approach with the existing measures and show that it provides a distinctive, novel measure of rationality. Importantly, note that, whereas all these measures pertain to the analysis of the deterministic rational model, the maximal separation approach is applicable to any possible rational or nonrational, deterministic or stochastic model of choice.

Recently, Liang (2019) has explored whether the inconsistency part of a choice data set can be attributed to choice error or preference heterogeneity. More concretely, Liang adopts the flexible multiple-preference framework of Kalai, Rubinstein, and Spiegler (2002), in which the individual can use different preferences in different menus.[4] Liang (2019) envisions inconsistencies as being driven by two different mechanisms: (i) preference heterogeneity, represented by a large fraction of choices being explained by an, ideally small, set of the individual's preferences á la Kalai, Rubinstein, and Spiegler (2002); and (ii) error, represented by a small fraction of choices being captured by preferences outside that set. Although we share with Liang (2019) an interest in identifying the part of the data that is due to error, our approach differs in two ways: firstly, by replacing the multiple-preference framework with a methodology that applies to any stochastic choice model; and, secondly, as discussed previously, by providing both a goodness-of-fit measure and the best description of behavior.

3. Maximal Separations

Let X be a nonempty finite set of alternatives. Menus are nonempty subsets of alternatives, and, in order to accommodate diverse settings, such as consumer-type domains or laboratory-type domains, we consider a nonempty arbitrary domain of menus $\mathcal{D}$. Pairs (a, A), with $a \in A$ and $A \in \mathcal{D}$, are called observations, and denoted by $\mathcal{O}$. A stochastic choice function is a mapping $\sigma : \mathcal{O} \to [0, 1]$, which, for every $A \in \mathcal{D}$, satisfies that $\sum_{a \in A} \sigma(a, A) = 1$. We interpret $\sigma(a, A)$ as the probability of choosing alternative a in menu A. We denote by SCF the space of all stochastic choice functions. The data are represented by means of a stochastic choice function, which we denote by

---

3. Other influential approaches provide only a goodness-of-fit measure; these include Swofford and Whitney (1987), Famulari (1995), and Echenique, Lee, and Shum (2011), whose proposal is to focus on the number of violations of a rationality axiom, for example, weak axiom of revealed preference, contained in the data.

4. Crawford and Pendakur (2013) implement the approach of Kalai, Rubinstein, and Spiegler (2002) using a set of data on milk purchases, finding that five preferences are enough to fully rationalize the data. Apesteguia and Ballester (2010) study the computational complexity of finding the minimal number of multiple preferences that rationalize the data.

ρ and which we assume to be within SCF.[5] That is, $\rho(a, A) > 0$ for every $(a, A) \in \mathcal{O}$. A model is a nonempty closed subset Δ of SCF, representing all the possible stochastic choice functions consistent with the entertained theoretical model. We emphasize that, other than the considered model being closed, we make no further restrictions. Thus, the model Δ encompasses all the relevant randomness considered by the analyst. This may include a base theoretical model, and considerations on measurement error or unobserved heterogeneity. A model instance, that is, a particular member of the set of theoretically admissible stochastic choice functions, is typically denoted by $\delta \in \Delta$.

We say that $\langle \lambda, \delta, \varepsilon \rangle \in [0, 1] \times \Delta \times$ SCF is a separation of data ρ whenever $\rho = \lambda \delta + (1 - \lambda)\varepsilon$. In a separation, we write ρ as a convex combination of the stochastic choice function δ, which contains randomness consistent with model Δ, and the stochastic choice function ε, which represents unstructured residual behavior. The fraction of data explained by the model in the separation is given by the parameter λ. We are particularly interested in explaining the largest possible fraction of data using model Δ. We say that a separation $\langle \lambda^*, \delta^*, \varepsilon^* \rangle$ is maximal if there exists no other separation $\langle \lambda, \delta, \varepsilon \rangle$ with $\lambda > \lambda^*$. The following proposition shows the existence of maximal separations and facilitates their computation.[6]

PROPOSITION 1. *Maximal separations always exist and are characterized by:*

1. $\lambda^* = \max\limits_{\delta \in \Delta} \min\limits_{(a,A) \in \mathcal{O}} \frac{\rho(a,A)}{\delta(a,A)}$;

2. $\delta^* \in \arg\max\limits_{\delta \in \Delta} \min\limits_{(a,A) \in \mathcal{O}} \frac{\rho(a,A)}{\delta(a,A)}$; *and*

3. $\varepsilon^* = \frac{\rho - \lambda^* \delta^*}{1 - \lambda^*}$.

In order to grasp the logic implicit in Proposition 1, let us consider the nontrivial case where $\rho \notin \Delta$. Consider any model instance $\delta \in \Delta$. Then, for $\langle \lambda, \delta, \varepsilon \rangle$ to be a separation of ρ, the residual stochastic choice function ε must lie on the line defined by ρ and δ, with ρ in between δ and ε. Now, notice that we can always trivially consider the separation $\langle 0, \delta, \rho \rangle$, where all data are regarded as residual behavior. To obtain larger values of λ with instance δ, ε must deviate from ρ in the opposite direction to that taken by δ. Ultimately, λ will be maximal when the residual behavior ε reaches the frontier of SCF, that is, when some observation has probability 0 or 1. Indeed, we only need to consider the case $\varepsilon(a, A) = 0$, that is, $\rho(a, A) < \delta(a, A)$ or, equivalently, $\rho(a, A)/\delta(a, A) < 1$, because, if $\varepsilon(a, A) = 1$ for some observation, we must also have that $\varepsilon(b, A) = 0$ for any other alternative $b \in A \backslash \{a\}$. Trivially, $\varepsilon(a, A) = 0$ is equivalent to $\lambda = \rho(a, A)/\delta(a, A)$, and, hence, the frontier will be first reached by the observation

---

5. This assumption is for expositional convenience; the case of ρ in the boundary of SCF can be dealt with trivially.

6. In order to avoid the discussion of indeterminacy in fractions throughout the text, we set the ratio $\rho(a, A)/0$ to be strictly larger than any real number. This is a harmless convention, because we could simply replace the expression $\min_{(a,A) \in \mathcal{O}}(\rho(a, A)/\delta(a, A))$ with $\min_{(a,A) \in \mathcal{O}, \delta(a,A) \neq 0}(\rho(a, A)/\delta(a, A))$. Moreover, whenever $\lambda^* = 1$, ε^* is any stochastic choice function. All proofs are given in Appendix A.

that minimizes the ratio $\rho(a, A)/\delta(a, A)$. Because these observations will play a key role in our analysis, we provide a formal definition here. Given instance δ, the set of observations that minimize the ratio $\rho(a, A)/\delta(a, A)$ are called δ-critical observations and are denoted by $\mathcal{O}_\delta$. Obviously, the maximal fraction of data that can be explained with instance δ is $\min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta(a, A))$, or, equivalently, $\rho(a, A)/\delta(a, A)$ for any $(a, A) \in \mathcal{O}_\delta$. When all possible instances of the model Δ are considered, the result follows.

As already mentioned, Proposition 1 works for arbitrary domains of menus. One domain, which has received a great deal of attention in the stochastic choice literature, is that of binary menus. Because we will also be using this domain in our experimental application, it is worth mentioning that it is one in which Proposition 1 is particularly simple to apply. In essence, notice that any model instance will overpredict the probability of choice of one of the alternatives in each binary menu within the domain, while underpredicting the other. Thus, one instance of the model is able to explain a fraction of the data that can be computed by looking at the least overpredicted alternative among all pairs.

4. Particular Models of Choice

Section 3 characterizes maximal separations for every possible model Δ. We now work with specific choice models. In each case, we use Proposition 1, together with the particular structure of the model under investigation, to offer more targeted results on maximal separations. The models we consider are the deterministic choice model and three stochastic choice models incorporating different forms of randomness: the tremble model, the Luce model, and the SCRUM. The three stochastic choice models have the deterministic model as a special case, but are mutually independent. Appendix B illustrates the application of each of the results developed here using a simple example involving three alternatives.

4.1. Deterministic Rationality

The standard economic decision-making model contemplates no randomness whatsoever. Behavior is deterministic and described as the outcome of the maximization of a single preference relation. Thus, in light of the deterministic model, all behavioral randomness must be regarded as residual behavior. Formally, denote by $\mathcal{P}$ the collection of all strict preference relations, that is, all transitive, complete, and asymmetric binary relations on X. Maximization of $P \in \mathcal{P}$ generates the deterministic rational choice function δ_P, which assigns probability 1 to the maximal alternative in menu A according to preference P. We denote this alternative by $m_P(A)$, that is, $m_P(A) \in A$ and $m_P(A)Py$ for every $y \in A \backslash \{m_P(A)\}$. Denote by DET the model composed of all the deterministic rational choice functions.

The following result shows that the maximal separation for DET can be easily computed using a simple recursive structure on subdomains of the data. For presenting

the result, some notation will be useful. Given a subset $S \subseteq X$, denote by $\mathcal{D}|_S = \{A \in \mathcal{D} : A \subseteq S\}$ and $\mathcal{O}|_S = \{(a, A) \in \mathcal{O} : A \subseteq S\}$ the corresponding subdomains of menus and observations involving subsets of S. Then, we have the following proposition.

PROPOSITION 2. *Let $\{\lambda_S\}_{S:\mathcal{D}|_S \neq \varnothing}$ and $P \in \mathcal{P}$ satisfy:*

1. $\lambda_S = \max\limits_{a \in S} \min \left\{ \{\rho(a, A)\}_{(a,A) \in \mathcal{O}|_S}, \lambda_{S \setminus \{a\}} \right\}$;

2. $m_P(S) \in \arg \max\limits_{a \in S} \min \left\{ \{\rho(a, A)\}_{(a,A) \in \mathcal{O}|_S}, \lambda_{S \setminus \{a\}} \right\}$. [7]

Then, $\langle \lambda_X, \delta_P, (\rho - \lambda_X \delta_P)/(1 - \lambda_X) \rangle$ is a maximal separation for the deterministic model.

Proposition 2 enables a recursive computation of maximal separations for DET. More precisely, the algorithm constructs a maximal separation for each restriction of data ρ to a subdomain of menus $\mathcal{D}|_S$, starting with subdomains in which $\mathcal{D}|_S = \{S\}$, that is, menus for which there are no available data in proper subsets. In these menus, only the highest choice frequency of an alternative must be considered. The maximal separation can be constructed by considering the preference relation that ranks the alternative with the highest choice frequency above all other alternatives. For any other subdomain $\mathcal{D}|_S$, the algorithm must analyze the alternatives $a \in S$ one by one, again considering the consequences of placing a as the maximal alternative in S. It turns out that we only need to consider the following values: (i) the choice frequencies of a in subsets of S, and (ii) the previously computed maximal fractions over the subdomains where alternative a is not present.[8]

4.2. Tremble Model

In tremble models, behavioral randomness is interpreted as a choice error. In the simplest version, the individual contemplates a preference relation P. With probability $(1 - \gamma) \in [0, 1]$, the preference is maximized. With probability γ, the individual trembles and randomizes among all the alternatives in the menu.[9] This generates the tremble choice function $\delta_{[P, \gamma]}(a, A) = \gamma/|A|$ whenever $a \in A \setminus \{m_P(A)\}$ and $\delta_{[P, \gamma]}(m_P(A), A) = 1 - \gamma(|A| - 1)/|A|$. Denote by Tremble the model composed of all tremble choice functions. The following result describes the maximal fraction of data explained by Tremble and a maximal separation for Tremble.

---

7. Notice that equations (1) and (2) always compute a minimum over a nonempty collection of values. This is because the computation takes place only when $\mathcal{D}|_S$ is nonempty, and, hence, either $a \in A$ for some $A \subseteq S$, or $\mathcal{D}|_{S \setminus \{a\}} \neq \varnothing$.

8. A particularly interesting example involves binary domains in which some stochastic transitivity property is satisfied. In this case, it is easy to see that the identified preference will be consistent with the stochastic revealed preference.

9. See Harless and Camerer (1994) for an early treatment of the trembling-hand concept in the stochastic choice literature.

PROPOSITION 3.   *Let $\{\lambda_S(\gamma)\}_{S:\mathcal{D}|_S \neq \varnothing}$ and $P(\gamma) \in \mathcal{P}$ satisfy, for every $\gamma \in [0, 1]$:*

1. $\lambda_S(\gamma) = \max\limits_{a \in S} \min \left\{ \left\{ \dfrac{|A|\rho(a,A)}{(1-\gamma)|A|+\gamma} \right\}_{(a,A)\in\mathcal{O}|_S}, \left\{ \dfrac{|A|\rho(b,A)}{\gamma} \right\}_{\substack{(b,A)\in\mathcal{O}|_S \\ b\neq a\in A}}, \lambda_{S\setminus\{a\}}(\gamma) \right\};$

2. $m_{P(\gamma)}(S) \in \arg\max\limits_{a \in S} \min \left\{ \left\{ \dfrac{|A|\rho(a,A)}{(1-\gamma)|A|+\gamma} \right\}_{(a,A)\in\mathcal{O}|_S}, \right.$

$$\left. \left\{ \dfrac{|A|\rho(b,A)}{\gamma} \right\}_{\substack{(b,A)\in\mathcal{O}|_S \\ b\neq a\in A}}, \lambda_{S\setminus\{a\}}(\gamma) \right\}.$$

Let γ^* be the tremble value that maximizes $\lambda_X(\gamma)$. Then,

$$\left\langle \lambda_X(\gamma^*), \delta_{[P(\gamma^*),\gamma^*]}, \frac{\rho - \lambda_X(\gamma^*)\delta_{[P(\gamma^*),\gamma^*]}}{1 - \lambda_X(\gamma^*)} \right\rangle$$

is a maximal separation for the tremble model.

Given the immediate connection with the rational deterministic model, the intuition of the result is analogous to that in Proposition 2.[10]

4.3. Luce Model

Denote by $\mathcal{U}$ the collection of strictly positive utility functions u such that, without loss of generality, $\sum_{x\in X} u(x) = 1$. Given $u \in \mathcal{U}$, a strictly positive Luce stochastic choice function is defined by $\delta_u(a, A) = u(a)/(\sum_{b\in A} u(b))$ with $a \in A \in \mathcal{D}$. In order to accommodate the Luce model within our framework, we consider the closure of the set of strictly positive Luce stochastic choice functions, which we denote by Luce.[11] We write δ_L to denote a generic, not necessarily strictly positive, Luce stochastic choice function. However, as shown in the proof of Proposition 4, strictly positive instances of the Luce model are always identified in maximal separations, and, hence, the previous assumption is inconsequential.

We now describe the structure of maximal separations of Luce. From Proposition 1, we know that the study of a particular instance of model δ_L requires us to analyze its critical observations $\mathcal{O}_{\delta_L}$. It turns out that, under the Luce model, we only need to check for a simple condition on the set $\mathcal{O}_{\delta_L}$.

---

10.   As in the deterministic case with binary domains, where choice satisfies stochastic transitivity, the maximal separation for the tremble model identifies the preference relation that is consistent with the stochastic revealed preference. Hence, in this case, both the deterministic and the tremble models identify the same preference relation. Interestingly, this is precisely the case in our empirical application. However, as we show in Appendix B, generally speaking, the maximal separations for the deterministic and tremble models do not necessarily identify the same preference relation.

11.   Effectively, the added stochastic choice functions have zero choice probabilities in some observations, and Luce-type behavior otherwise. See Echenique and Saito (2019) and Horan (2019) for studies of the treatment of zero choice probabilities in models à la Luce.

PROPOSITION 4. $\left\langle \min_{(a,A)} \frac{\rho(a,A)}{\delta_L^*(a,A)}, \delta_L^*, \frac{\rho-\lambda^*\delta_L^*}{1-\lambda^*} \right\rangle$ *is a maximal separation for the Luce model if and only if* $\mathcal{O}_{\delta_L^*}$ *contains a subcollection* $\{(a_i, A_i)\}_{i=1}^I$ *such that* $\bigcup_{i=1}^I \{a_i\} = \bigcup_{i=1}^I A_i$.

Proposition 4 provides a simple means to obtain maximal separations for the Luce model, which entails checking whether the critical observations of a Luce stochastic choice function satisfy a cyclical property. Consider a strictly positive instance of Luce given by $u \in \mathcal{U}$ and its critical observations $\mathcal{O}_{\delta_u}$. Clearly, for another separation using a different Luce vector $v \in \mathcal{U}$ to explain a larger fraction of the data, it should be possible to improve critical observation (a_1, A_1) by reducing the predicted choice probability of a_1. This requires that one alternative in A_1, say, a_2, is such that $v(a_2)/v(a_1) > u(a_2)/u(a_1)$. However, because there exists a critical observation of the form (a_2, A_2), we need to find another alternative in A_2, say, a_3, with $v(a_3)/v(a_2) > u(a_3)/u(a_2)$. Given that $\bigcup_{i=1}^I \{a_i\} = \bigcup_{i=1}^I A_i$, this process leads to a cycle, and, consequently, the ρ/δ ratio of all the critical observations of δ_u cannot be improved, and is, therefore, optimal. The situation is entirely different when there is $x \in \bigcup_{i=1}^I A_i \setminus \bigcup_{i=1}^I \{a_i\}$, suggesting that the following simple algorithm identifies a maximal separation for the Luce model. Start with any vector of weights $u \in \mathcal{U}$. Take any $x \in \bigcup_{i=1}^I A_i \setminus \bigcup_{i=1}^I \{a_i\}$ and move the utilities along the segment $\alpha 1_x + (1-\alpha)u$, where 1_x is a function assigning a value 1 to x and a value 0 to any other alternative. Eventually, this leads to a new Luce vector, which explains a strictly larger fraction of the data. This ascending algorithm yields the maximal separation.

4.4. Single-Crossing Random Utility Model

In RUMs, there exists a probability distribution μ over the set of all possible preferences $\mathcal{P}$. At the choice stage, a preference is realized according to μ, and maximized, thereby determining the choice probabilities $\delta_\mu(a, A) = \sum_{P \in \mathcal{P}: a = m_P(A)} \mu(P)$, for every $(a, A) \in \mathcal{O}$. In other words, the choice probability of a given alternative within a menu is given by the sum of the probability masses associated with the preferences where the alternative is maximal within the menu.

The literature has often considered these models complex to work with, and offered more easily applicable models in restricted domains. Here, we focus on SCRUMs, which are RUMs over a set of preferences satisfying the single-crossing condition.[12] Formally, SCRUMs consider probability distributions μ on a given ordered collection of preferences $\mathcal{P}' = \{P_1, P_2, \ldots, P_T\}$, satisfying the single-crossing condition $P_j \cap P_1 \subseteq P_i \cap P_1$ if and only if $j \geq i$. That is, the preference over a pair of alternatives x and y reverses once at most in the ordered collection of preferences. We denote the set of

---

12. See Apesteguia, Ballester, and Lu (2017) for a study of this model. Other RUMs using restricted domains are Gul and Pesendorfer (2006) and Lu and Saito (2018).

TABLE 1. Lotteries.

$l_1 = (17)$	$l_4 = (30, 10)$	$l_7 = (40, 12, 5)$
$l_2 = (50, 0)$	$l_5 = (20, 15)$	$l_8 = (30, 12, 10)$
$l_3 = (40, 5)$	$l_6 = (50, 12, 0)$	$l_9 = (20, 12, 15)$

SCRUM stochastic choice functions by SC. Proposition 5 characterizes the maximal separations for SCRUMs.

PROPOSITION 5. *Let* $\lambda_1 = \min\limits_{A \in \mathcal{D}} \rho(m_{P_1}(A), A)$ *and* $\delta_{\mu_1} = \delta_{P_1}$, *and for every* $i \in \{2, \ldots, T\}$ *define recursively:*

1. $\lambda_i = \min\limits_{A \in \mathcal{D}} \left\{ \rho(m_{P_i}(A), A) + \max\limits_{j : j < i, m_{P_j}(A) \neq m_{P_i}(A)} \lambda_j \right\}$;

2. $\delta_{\mu_i} = \left(1 - \frac{\lambda_{i-1}}{\lambda_i}\right) \delta_{P_i} + \frac{\lambda_{i-1}}{\lambda_i} \delta_{\mu_{i-1}}$.

Then, $\left\langle \lambda_T, \delta_{\mu_T}, \frac{\rho - \lambda_T \delta_{\mu_T}}{1 - \lambda_T} \right\rangle$ *is a maximal separation for SCRUM.*

Proposition 5 provides a smooth recursive method with which to obtain a maximal separation. It basically computes the maximal fraction of data, λ_i, that can be explained by SCRUMs using preferences up to P_i. Trivially, the maximal fraction of data explained by P_1 is $\min\limits_{A \in \mathcal{D}} \rho(m_{P_1}(A), A)$. Now consider any other preference $P_i \in \mathcal{P}'$ and assume that every preference $P_j, j < i$, has been analyzed. With the extra preference P_i, and for a given menu A, we can rationalize data $\rho(m_{P_i}(A), A)$ together with any other data $\rho(x, A)$, $x \neq m_{P_i}(A)$, which are rationalized by preferences preceding P_i. This can be achieved by considering the appropriate linear combination of the constructed SCRUM that uses preferences up to P_{i-1} with preference P_i.

5. An Empirical Application

Here we use an experimental data set to operationalize the maximal separation results obtained in the previous section.[13] There are nine equiprobable monetary lotteries, described in Table 1. Each of the 87 participants faced 108 different menus of lotteries, including all 36 binary menus and a random sample of larger menus.[14] There are two

---

13.   We collected the experimental data together with Syngjoo Choi at UCL in March 2013, within the context of another research project. This is the first completed paper to use this data set. We are very grateful to Syngjoo for kindly allowing us to use it.

14.   Menus of two, three, and five alternatives were presented one at a time, in a randomized order. No participant was presented more than once with the same menu of alternatives. The location of the lotteries on the screen was randomized, as was the location of the monetary prizes within a lottery.

treatments. Treatment NTL is a standard implementation, with no time limit on the choice. In treatment TL, subjects had to select a lottery within a limited time. At the end of the experiment, one of the menus was chosen at random and the subject was paid according to his or her choice from that menu.[15]

To ensure a sufficiently large number of data points per menu, we focus on the choices made in the binary menus, which, when both treatments are aggregated, gives a total of 87 data points per menu.[16] Table 2 reports the choice probabilities in each of the binary menus. It also reports the optimal and residual stochastic choice functions identified in the maximal separation results, using the models described in the previous section. In SCRUM, we use the CRRA expected utility representation, which is by far the most widely used utility representation for risk preferences.[17] There are several lessons to be learned from the table.

First, note that the maximal fractions of the data explained by the respective models are successively increasing from the deterministic choice model to the tremble model, to the Luce model, and, finally, to the SCRUM-CRRA model. It is worth noting that the deterministic model already explains about half of the data, that is, 0.51.[18] The identified optimal instance is the one associated with the preference $l_1 P l_5 P l_4 P l_8 P l_7 P l_9 P l_3 P l_6 P l_2$. The top alternative, lottery l_1, is the safest, because it gives £17 with probability 1. The next is lottery l_5, which has the second lowest variance at the expense of a very low expected return. Lottery l_2, the one with the highest expected value and highest variance, is regarded as the worst alternative. Hence, the deterministic model depicts a population that is essentially highly risk-averse. The model reaches its explanatory limits with the critical observation $(l_8, \{l_7, l_8\})$, where, by Proposition 1, the ratio of observed to predicted probability is minimal. Specifically, the observed choice probability is 0.51, whereas the deterministic prediction is 1. The ratio of these two values gives the fraction of data explained by the model, 0.51.

The tremble model identifies exactly the same preference as the deterministic model, while increasing the maximal fraction of the data explained from 0.51 to 0.68. This is the result of using a relatively large tremble probability, $\gamma = 0.51$. The tremble model is characterized by critical observations $(l_8, \{l_7, l_8\})$ and $(l_9, \{l_5, l_9\})$. As in the

---

15. Specifically, subjects had 5, 7 and 9 seconds to choose from the menus of two, three, and five alternatives, respectively.

16. Due to the time limit in one of the treatments, the number is slightly lower for some menus. Specifically, there are 18 menus with 87 data points, 12 with 86, 3 with 85, and 3 with 84.

17. The CRRA Bernoulli function is $x^{1-r}/(1-r)$, whenever $r \neq 1$, and $\log x$ otherwise, with x representing money. We have also studied the cases of constant absolute risk aversion (CARA) expected utility, and mean-variance utility, and obtained similar results, which are available upon request. Note that SCRUM with CRRA is but a generalization of the random parameter model used in Apesteguia and Ballester (2018), in the sense that the former imposes no probability distribution over the set of preferences.

18. In order to put this result into perspective, consider Crawford and Pendakur's (2013) analysis of data from a Danish household survey on the purchase of six different types of milk. They find that a single preference relation is sufficient to rationalize 64% of the data. The Houtman–Maks index gives a consistency level of 66%. In Appendix D, we review this index, arguing that it is slightly more flexible than applying the maximal separation technique to the deterministic model, which explains the higher consistency found in the data.

TABLE 2. Data and maximal separations.

(a, A)	ρ	DET		TREMBLE		LUCE		SCRUM-CRRA	
		δ^*_{DET}	$\varepsilon^*_{\text{DET}}$	$\delta^*_{\text{Tremble}}$	$\varepsilon^*_{\text{Tremble}}$	δ^*_{Luce}	$\varepsilon^*_{\text{Luce}}$	$\delta^*_{\text{SC-CRRA}}$	$\varepsilon^*_{\text{SC-CRRA}}$
$(l_1, \{l_1, l_2\})$	0.75	1.00	0.49	0.74	0.75	0.91	0.30	0.74	0.77
$(l_1, \{l_1, l_3\})$	0.60	1.00	0.19	0.74	0.29	0.71	0.28	0.55	0.78
$(l_2, \{l_2, l_3\})$	0.33	0.00	0.67	0.26	0.50	0.20	0.69	0.24	0.66
$(l_1, \{l_1, l_4\})$	0.53	1.00	0.05	0.74	0.07	0.62	0.27	0.47	0.75
$(l_2, \{l_2, l_4\})$	0.28	0.00	0.56	0.26	0.32	0.15	0.64	0.24	0.40
$(l_3, \{l_3, l_4\})$	0.43	0.00	0.86	0.26	0.78	0.40	0.50	0.42	0.46
$(l_1, \{l_1, l_5\})$	0.58	1.00	0.16	0.74	0.24	0.46	0.92	**0.47**	**1.00**
$(l_2, \{l_2, l_5\})$	0.25	0.00	0.51	0.26	0.25	0.08	0.73	0.26	0.23
$(l_3, \{l_3, l_5\})$	0.45	0.00	0.92	0.26	0.87	**0.26**	**1.00**	0.45	0.46
$(l_4, \{l_4, l_5\})$	0.49	0.00	0.99	0.26	0.98	0.34	0.89	0.53	0.33
$(l_1, \{l_1, l_6\})$	0.72	1.00	0.44	0.74	0.68	0.87	0.31	0.76	0.60
$(l_2, \{l_2, l_6\})$	0.44	0.00	0.89	0.26	0.84	0.42	0.51	0.42	0.53
$(l_3, \{l_3, l_6\})$	0.80	1.00	0.60	0.74	0.93	**0.74**	**1.00**	0.79	0.84
$(l_4, \{l_4, l_6\})$	0.76	1.00	0.51	0.74	0.79	0.81	0.62	0.76	0.76
$(l_5, \{l_5, l_6\})$	0.75	1.00	0.49	0.74	0.75	0.89	0.35	0.76	0.71
$(l_1, \{l_1, l_7\})$	0.63	1.00	0.25	0.74	0.38	0.77	0.23	0.74	0.22
$(l_2, \{l_2, l_7\})$	0.24	0.00	0.49	0.26	0.22	0.26	0.21	0.26	0.19
$(l_3, \{l_3, l_7\})$	0.48	0.00	0.96	0.26	0.94	0.57	0.20	0.53	0.27
$(l_4, \{l_4, l_7\})$	0.62	1.00	0.24	0.74	0.37	0.67	0.49	0.76	0.14
$(l_5, \{l_5, l_7\})$	0.63	1.00	0.26	0.74	0.40	0.79	0.18	0.76	0.18
$(l_6, \{l_6, l_7\})$	0.27	0.00	0.54	0.26	0.29	0.33	0.10	0.24	0.36
$(l_1, \{l_1, l_8\})$	0.64	1.00	0.27	0.74	0.42	0.67	0.57	0.76	0.21
$(l_2, \{l_2, l_8\})$	0.22	0.00	0.45	0.26	0.15	0.17	0.36	0.26	0.09
$(l_3, \{l_3, l_8\})$	0.36	0.00	0.73	0.26	0.58	0.45	0.12	0.45	0.03
$(l_4, \{l_4, l_8\})$	0.56	1.00	0.12	0.74	0.18	0.55	0.60	0.56	0.56
$(l_5, \{l_5, l_8\})$	0.62	1.00	0.23	0.74	0.36	0.70	0.40	0.76	0.13
$(l_6, \{l_6, l_8\})$	0.20	0.00	0.40	0.26	0.07	0.23	0.12	0.24	0.04
$(l_7, \{l_7, l_8\})$	0.49	**0.00**	**1.00**	**0.26**	**1.00**	0.37	0.83	0.42	0.77
$(l_1, \{l_1, l_9\})$	0.76	1.00	0.51	0.74	0.78	0.74	0.81	0.79	0.62
$(l_2, \{l_2, l_9\})$	0.28	0.00	0.56	0.26	0.32	0.23	0.42	0.28	0.28
$(l_3, \{l_3, l_9\})$	0.39	0.00	0.79	0.26	0.68	**0.53**	**0.00**	0.45	0.17
$(l_4, \{l_4, l_9\})$	0.55	1.00	0.08	0.74	0.13	0.63	0.32	0.53	0.60
$(l_5, \{l_5, l_9\})$	0.83	1.00	0.65	**0.74**	**1.00**	**0.76**	**1.00**	1.00	0.20
$(l_6, \{l_6, l_9\})$	0.22	0.00	0.44	0.26	0.14	0.29	0.02	0.26	0.08
$(l_7, \{l_7, l_9\})$	0.56	1.00	0.12	0.74	0.18	0.46	0.87	0.45	0.96
$(l_8, \{l_8, l_9\})$	0.64	1.00	0.26	0.74	0.41	0.58	0.78	**0.53**	**1.00**
λ^*_Δ		0.51		0.68		0.74		0.78	

Notes: (a, A) denotes the observation referring to alternative a from menu A, ρ the observed frequency of choosing lottery a from menu A, and $\langle \lambda^*_\Delta, \delta^*_\Delta, \varepsilon^*_\Delta \rangle$ the maximal separation of ρ for model $\Delta \in$ {DET, Tremble, Luce, SC-CRRA}. Data entries in bold refer to the menus containing the critical observations in the respective model.

deterministic case, choice data are scarce for l_8 versus l_7, but the problem is less severe thanks to the presence of a tremble, due to which the individual is predicted to choose l_8 with a probability of only 0.74, which reduces the ratio of observed to predicted probabilities to 0.68. This ratio cannot be improved beyond this point. Although a higher tremble probability would increase this ratio, it would also decrease the ratio of the other critical observation, $(l_9, \{l_5, l_9\})$, which has the same value of 0.68. To see this, notice that the choice prediction for alternative l_9, being worse than alternative l_5, corresponds entirely to the tremble probability, and, hence, an increase in tremble would increase the predicted probability and, thus, decrease the ratio.

The Luce model is able to explain close to three quarters of the data. The optimal Luce utility values suggest a highly risk-averse population. Although u does not represent P_{DET} exactly, it represents a preference very close to it. Interestingly, we see that the Luce model can accommodate a larger fraction of the data by allowing randomness to depend on the cardinal evaluation of the alternatives. The model is hard pressed to explain observations $(l_5, \{l_3, l_5\})$, $(l_6, \{l_3, l_6\})$, $(l_3, \{l_3, l_9\})$ and $(l_9, \{l_5, l_9\})$, which have the type of cyclical structure described in Proposition 4. In each of these observations, the ratio of observed to predicted probabilities is equal to 0.74, which is the Luce critical value.

Finally, SC-CRRA explains 78% of the data. In so doing, it assigns positive masses to 10 of the 30 possible CRRA preferences; the largest probability mass, 0.44, is assigned to the most risk-averse CRRA preference, that is, preference $l_1 P l_5 P l_9 P l_8 P l_4 P l_7 P l_3 P l_6 P l_2$, which is again very close to P_{DET}. Because each preference compatible with CRRA corresponds to an interval of risk-aversion levels, we can completely describe the optimal SC-CRRA instance by reporting the values of the cumulative distribution function at the upper bounds of these intervals. These are $F(-4.15) = 0.205$, $F(-0.31) = 0.241$, $F(-0.08) = 0.242$, $F(0.34) = 0.258$, $F(0.41) = 0.276$, $F(0.44) = 0.416$, $F(0.61) = 0.453$, $F(1) = 0.533$, $F(4.71) = 0.563$, and $F(\infty) = 1$. Notice that, in addition to explaining a large fraction of the data, SC-CRRA is also rich enough to show that a quarter of the population is risk loving, $F(-0.08) = 0.242$. SC-CRRA reaches the limits of its explanatory power at observations $(l_5, \{l_1, l_5\})$ and $(l_9, \{l_8, l_9\})$. On one hand, lottery l_5 is preferred over lottery l_1 by all CRRA levels below 2, which has an accumulated mass of 0.533. Given the observed choices, this leads to a critical ratio for observation $(l_5, \{l_1, l_5\})$ of 0.78. Improving this ratio would necessarily require us to assign a higher weight to risk-aversion levels higher than 2. However, this would immediately conflict with the ratio of l_9 to l_8, because l_9 is ranked above l_8 at all risk-aversion levels higher than 1. As the ratio of observed to predicted data for $(l_9, \{l_8, l_9\})$ also has the critical value of 0.78, no further improvement is possible.[19]

To conclude the discussion of Table 2, we would like to emphasize that the four models are very consistent in their qualitative descriptions of the population, all judging

---

19. In Appendix C, we use this data set to analyze the maximal separation using Gul and Pesendorfer's (2006) random expected utility (REU) model.

it to be highly risk averse. We then see that, by introducing different sources of randomness, it is possible to explain larger fractions of the data, and that the precise source of randomness affects the fraction of the data explained.

6. Other Goodness-of-Fit Measures

The maximal separation exercise identifies a best instance of the model $\delta^* \in \Delta$ and an expression of residual behavior $\varepsilon^* \in \text{SCF}$, which, combined at rates λ^* and $1 - \lambda^*$, generate data ρ. The value λ^* is a tight upper bound for the fraction of data that can be explained by the model. Thus, the exercise provides a measure of the goodness of fit of model Δ to data ρ. There are other well-known measures in the literature that partially share the structure of the maximal separation measure, in the sense that they also identify one instance of the model that maximizes a notion of closeness to the data.[20] For the sake of comparison, we adopt the standard language of minimization of loss functions in speaking of lack of fit all throughout the section.[21]

Formally, a loss function is a map $L : \Delta \times \rho \rightarrow \mathbb{R}_+$ that measures the deviation of every instance $\delta \in \Delta$ with respect to data ρ. The lack of fit and the best instance of the model follow immediately from the minimization of the loss function among the different instances of the model.[22] Now, in a maximal separation, the minimal fraction of data unexplained, $1 - \lambda^*$, represents a measure of the lack of fit, which can be written as the minimization of a loss function. From Proposition 1, we know that $1 - \lambda^* = 1 - \max_{\delta \in \Delta} \min_{(a,A) \in \mathcal{O}} (\rho(a, A)/\delta(a, A)) = \min_{\delta \in \Delta} [\max_{(a,A) \in \mathcal{O}} (1 - \rho(a, A)/\delta(a, A))]$, and, hence, we can write the maximal separation loss function as $L_{\text{MS}}(\delta, \rho) = \max_{(a,A) \in \mathcal{O}} [1 - \rho(a, A)/\delta(a, A)]$.

Two other important goodness-of-fit measures are maximum likelihood and least squares. The maximum likelihood exercise entails the minimization of the Kullback–Leibler divergence from δ to ρ, which can be written as the minimization of the loss function $L_{\text{ML}}(\delta, \rho) = \sum_{(a,A) \in \mathcal{O}} \rho(a, A) \log(\rho(a, A)/\delta(a, A))$.[23] Similarly, least squares entails the minimization of the quadratic loss function $L_{\text{LS}}(\delta, \rho) = \sum_{(a,A) \in \mathcal{O}} (\delta(a, A) - \rho(a, A))^2$.

---

20.   We say that the other measures partially share the structure of maximal separations because they do not identify (minimal) expressions of residual behavior. This component, which we believe potentially crucial for the understanding of actual behavior and revision of theoretical models, is unique to maximal separations.

21.   By lack of fit, sometimes also known as badness of fit, we mean the mirror notion of goodness of fit, basically, how poorly a model fits the data.

22.   Notice that, when the model Δ is the deterministic rational model of choice, lack of fit merely corresponds to a notion of irrationality of the data. As mentioned in Section 2, most measures of irrationality, including Afriat, Varian, Houtman–Maks, and the Swaps Index, adopt this minimization structure. For the specific case of the deterministic model, Appendix D formally compares the maximal separation approach with other rationality measures.

23.   The Kullback–Leibler divergence can be interpreted as the amount of information lost due to the use of δ instead of ρ.

TABLE 3. Maximal separation and maximum likelihood.

Deterministic		
MS	$P = [l_1, l_5, l_4, l_8, l_7, l_9, l_3, l_6, l_2]$	
ML	$P = [l_1, l_5, l_4, l_8, l_7, l_9, l_3, l_6, l_2]$	
Tremble		
MS	$P = [l_1, l_5, l_4, l_8, l_7, l_9, l_3, l_6, l_2]$;	$\gamma = 0.51$
ML	$P = [l_1, l_5, l_4, l_8, l_7, l_9, l_3, l_6, l_2]$;	$\gamma = 0.68$
Luce		
MS	$u = (0.22, 0.02, 0.09, 0.13, 0.25, 0.03, 0.07, 0.11, 0.08)$	
ML	$u = (0.18, 0.04, 0.1, 0.14, 0.17, 0.04, 0.11, 0.13, 0.09)$	

SCRUM-CRRA

MS $F(-4.15) = 0.205, F(-0.31) = 0.241, F(-0.08) = 0.242, F(0.34) = 0.258, F(0.41) = 0.276$
$F(0.44) = 0.416, F(0.61) = 0.453, F(1) = 0.533, F(4.71) = 0.563, F(\infty) = 1$

ML $F(-4.15) = 0.22, F(-0.31) = 0.287, F(0.44) = 0.442$
$F(1) = 0.506, F(4.71) = 0.563, F(\infty) = 1$

Notes: MS and ML denote maximal separation and maximum likelihood, respectively. P denotes the preference identified in each respective case, where the ranking declines from the left- to right-hand side, γ is the tremble probability in Tremble, u is the Luce utility vector associated with Luce, where the ith entry in u corresponds to the utility value of lottery l_i, and, finally, $F(r)$ denotes the cumulative probability masses associated with the upper bounds of the intervals of the relative risk-aversion coefficients r consistent with those CRRA preference relations that have a strictly positive mass in the corresponding estimation procedure.

On inspecting the loss functions, it becomes immediately clear that the maximal separation measure is different from those defined by maximum likelihood and least squares. Crucially, whereas the maximal separation is concerned with the largest deviation between the data and the specified model, maximum likelihood and least squares aggregate the deviations across the different observations. This has two implications. First, there should be data sets and models where maximal separation identifies different best instances of the model. Secondly, we should expect maximal separation to provide more accurate overestimations for those observations for which the observed choice frequency is low, whereas the other measures would perform better on average. In what follows, we use our experimental data set to illustrate these two points empirically.

Table 3 illustrates the first point. It reports the instances of the models identified by the maximal separation and the maximum likelihood techniques over the entire data set.[24] No difference whatsoever is observed with respect to the deterministic model; the estimated preference relations are exactly the same. This ordinal equivalence

---

24. In the ML calculations, we impose a lower bound in the theoretical predictions, in order to ensure strictly positive likelihoods. The results given by least squares are practically identical to those given by maximum likelihood, and are, therefore, omitted.

is preserved in the case of the tremble model, although our technique predicts a substantially smaller trembling coefficient, $0.51 < 0.68$. The intuition for this difference is straightforward. Recall that, as mentioned previously, $(l_9, \{l_5, l_9\})$ is a critical observation in the maximal separation exercise for Tremble. The observed probability in this observation is small, 0.17, and, due to the trembling parameter, the instance of the model identified by our technique predicts a rather large relative frequency of 0.26. However, the maximum likelihood exercise is not severely affected by this local consideration and makes the estimation by simply averaging over all the observations. Consequently, the estimation exercise in maximum likelihood is willing to sacrifice the predictive accuracy of this extreme observation in order to better accommodate the more moderate ones. This is done by substantially increasing the trembling parameter and, consequently, the prediction in this particular observation $(l_9, \{l_5, l_9\})$, which reaches a disproportionate value of 0.34, which is twice the observed value. Similar reasoning applies to the comparison of the Luce and SC-CRRA cases.

For a comparative illustration of the different approaches, we now perform an out-of-sample exercise. This will also allow us to evaluate the second conjecture stated previously.[25] We take all the binary data except for one binary set, estimate the instances of the models by maximal separation and maximum likelihood using these data, and use the estimated instances to predict the behavior in the omitted binary set. We perform this procedure on 36 binary sets, each with 2 cases: one in which both maximal separation and maximum likelihood overestimate the choice probability of the same alternative in the binary menu and another in which they overestimate the choice probability of different alternatives. By focusing on the first case, a comparison of the predictive powers of maximal separation and maximum likelihood becomes straightforward; one of the methods is unambiguously more accurate than the other.[26] We, therefore, focus our comparison on these menus, because the conclusions may otherwise depend on the choice of distance function. Table 4 reports the results.[27] As advanced previously, the analysis of the loss functions entailed by the two techniques suggested a very intuitive conjecture. Namely, that the maximal separation technique is very cautious and can, therefore, be expected to perform better in observations with low choice probabilities. This conjecture is largely confirmed in our analysis. In all three models, the overestimation of small probabilities is less problematic for the maximal separation technique, whereas maximum likelihood deals better with the overestimation of large probabilities. We conclude from these results, therefore, that

---

25.   We complement this exercise in Appendix C by using the nonbinary part of the data set, whereas, in Appendix B, we elaborate on the intuition behind this conjecture using a more theoretical approach entailing the application of a particular data-generating process and the tremble model.

26.   Notice that, in binary menus, if one alternative is overestimated, the other is underestimated, and, for both observations, there is one method that is more accurate than the other. Thus, there is no loss of generality in discussing the results for, say, the overestimated alternatives.

27.   We do not report the results of the deterministic method, because, in this case, the maximal separation and maximum likelihood predictions are exactly the same.

TABLE 4. Forecasting results of maximal separation and maximum likelihood.

Tremble				Luce				SCRUM-CRRA			
(a, A)	ρ	MS	ML	(a, A)	ρ	MS	ML	(a, A)	ρ	MS	ML
$(l_9, \{l_5, l_9\})$	0.17	**0.28**	0.35	$(l_9, \{l_5, l_9\})$	0.17	**0.39**	0.41	$(l_6, \{l_3, l_6\})$	0.20	**0.21**	0.27
$(l_6, \{l_3, l_6\})$	0.20	**0.26**	0.35	$(l_6, \{l_3, l_6\})$	0.20	**0.27**	0.32	$(l_6, \{l_6, l_8\})$	0.20	**0.26**	0.29
$(l_6, \{l_6, l_8\})$	0.20	**0.26**	0.35	$(l_6, \{l_6, l_8\})$	0.20	**0.23**	0.25	$(l_6, \{l_6, l_9\})$	0.22	**0.26**	0.29
$(l_6, \{l_6, l_9\})$	0.22	**0.26**	0.34	$(l_6, \{l_6, l_9\})$	0.22	**0.29**	0.34	$(l_2, \{l_2, l_8\})$	0.22	**0.26**	0.29
$(l_2, \{l_2, l_8\})$	0.22	**0.26**	0.34	$(l_2, \{l_2, l_7\})$	0.24	**0.26**	0.29	$(l_6, \{l_4, l_6\})$	0.24	**0.24**	0.29
$(l_6, \{l_4, l_6\})$	0.24	**0.26**	0.34	$(l_9, \{l_1, l_9\})$	0.24	**0.26**	0.36	$(l_2, \{l_2, l_7\})$	0.24	**0.26**	0.29
$(l_2, \{l_2, l_7\})$	0.24	**0.26**	0.34	$(l_3, \{l_3, l_8\})$	0.36	0.49	*0.46*	$(l_2, \{l_1, l_2\})$	0.25	**0.26**	0.29
$(l_9, \{l_1, l_9\})$	0.24	**0.26**	0.34	$(l_9, \{l_8, l_9\})$	0.36	**0.42**	0.43	$(l_2, \{l_2, l_5\})$	0.25	**0.26**	0.29
$(l_2, \{l_1, l_2\})$	0.25	**0.26**	0.34	$(l_3, \{l_3, l_9\})$	0.39	0.63	*0.56*	$(l_3, \{l_3, l_8\})$	0.36	0.47	*0.45*
$(l_2, \{l_2, l_5\})$	0.25	**0.26**	0.34	$(l_5, \{l_1, l_5\})$	0.42	0.58	*0.50*	$(l_9, \{l_8, l_9\})$	0.36	0.55	*0.52*
$(l_6, \{l_5, l_6\})$	0.25	**0.26**	0.34	$(l_9, \{l_7, l_9\})$	0.44	0.54	*0.46*	$(l_3, \{l_3, l_9\})$	0.39	**0.45**	0.48
$(l_3, \{l_3, l_9\})$	0.39	0.74	*0.66*	$(l_8, \{l_4, l_8\})$	0.44	**0.45**	0.49	$(l_3, \{l_1, l_3\})$	0.40	0.45	*0.45*
$(l_2, \{l_2, l_6\})$	0.44	0.74	*0.66*	$(l_8, \{l_7, l_8\})$	0.51	0.59	*0.55*	$(l_5, \{l_1, l_5\})$	0.42	0.56	*0.56*
$(l_4, \{l_4, l_5\})$	0.49	0.74	*0.66*	$(l_5, \{l_4, l_5\})$	0.51	0.68	*0.54*	$(l_9, \{l_7, l_9\})$	0.44	0.58	*0.58*
$(l_7, \{l_7, l_8\})$	0.49	0.75	*0.66*	$(l_1, \{l_1, l_4\})$	0.53	0.63	*0.56*	$(l_9, \{l_4, l_9\})$	0.45	**0.47**	0.50
$(l_7, \{l_3, l_7\})$	0.52	0.74	*0.66*	$(l_5, \{l_3, l_5\})$	0.55	0.78	*0.65*	$(l_4, \{l_1, l_4\})$	0.47	0.53	*0.51*
$(l_1, \{l_1, l_4\})$	0.53	0.74	*0.66*	$(l_4, \{l_4, l_9\})$	0.55	0.64	*0.63*	$(l_3, \{l_3, l_4\})$	0.48	0.53	*0.51*
$(l_5, \{l_3, l_5\})$	0.55	0.74	*0.66*	$(l_4, \{l_3, l_4\})$	0.57	0.60	*0.60*	$(l_4, \{l_4, l_5\})$	0.49	0.53	*0.51*
$(l_4, \{l_4, l_9\})$	0.55	0.74	*0.66*	$(l_1, \{l_1, l_3\})$	0.60	0.71	*0.66*	$(l_8, \{l_7, l_8\})$	0.51	0.61	*0.57*
$(l_7, \{l_7, l_9\})$	0.56	0.74	*0.66*	$(l_3, \{l_2, l_3\})$	0.67	0.80	*0.71*	$(l_5, \{l_3, l_5\})$	0.55	**0.55**	0.56
$(l_4, \{l_3, l_4\})$	0.57	0.74	*0.66*	$(l_4, \{l_2, l_4\})$	0.72	0.85	*0.78*	$(l_6, \{l_2, l_6\})$	0.56	0.58	*0.56*
$(l_1, \{l_1, l_3\})$	0.60	0.74	*0.66*	$(l_1, \{l_1, l_6\})$	0.72	0.87	*0.83*	$(l_5, \{l_5, l_8\})$	0.62	0.76	*0.72*
$(l_5, \{l_5, l_8\})$	0.62	0.74	*0.66*	$(l_1, \{l_1, l_2\})$	0.75	0.91	*0.82*	$(l_4, \{l_4, l_7\})$	0.62	0.79	*0.75*
$(l_4, \{l_4, l_7\})$	0.62	0.74	*0.66*	$(l_5, \{l_2, l_5\})$	0.75	0.92	*0.80*	$(l_1, \{l_1, l_7\})$	0.63	0.76	*0.72*
$(l_1, \{l_1, l_7\})$	0.63	0.74	*0.66*	$(l_5, \{l_5, l_6\})$	0.75	0.89	*0.81*	$(l_5, \{l_5, l_7\})$	0.63	0.76	*0.72*
$(l_5, \{l_5, l_7\})$	0.63	0.74	*0.66*	$(l_4, \{l_4, l_6\})$	0.76	0.83	*0.78*	$(l_1, \{l_1, l_8\})$	0.64	0.76	*0.72*
$(l_8, \{l_8, l_9\})$	0.64	0.74	*0.66*					$(l_3, \{l_2, l_3\})$	0.67	0.76	*0.72*
$(l_1, \{l_1, l_8\})$	0.64	0.74	*0.66*					$(l_1, \{l_1, l_9\})$	0.76	**0.84**	0.90
$(l_8, \{l_3, l_8\})$	0.64	0.74	*0.66*					$(l_5, \{l_5, l_9\})$	0.83	1.00	1.00

Notes: (a, A) denotes the observation referring to alternative a from menu A such that a is the lottery where the predictions of both maximal separation (MS) and maximum likelihood (ML) are above the observed choice data ρ. Those observations for which one of the predictions of MS or ML is above the observed choice data and the other below are not reported in the table. Then, for each one of the models, the binary menus of lotteries are ordered from lower to higher observed choice probabilities. Bold entries refer to the cases where MS is closer to the data and italicized entries refers to those cases where ML is closer to the data.

if the interest is in forecasting, it may be worth applying both maximal separation and maximum likelihood to obtain a clearer picture of the overall situation.

7. Discussion

We close this paper by commenting on three issues surrounding the notion of maximal separation. We begin by discussing how to select one of the available existing models by assigning a parsimony cost to each model. We then comment on the possibility of assuming that the model Δ is not only closed, but also convex. Finally, we discuss the possibility of restricting the space of residual stochastic choice functions, and comment on possible interpretations of residual behavior.

7.1. Model Selection

The fraction of data explained in a maximal separation constitutes an absolute performance measure, a concept in tension with the idea of overfitting, that is, larger models are explanatorily superior simply because of their size. For a direct example of this tension, notice that whenever $\Delta \subseteq \Delta'$, the maximal fraction of data explained by model Δ' is, independently of ρ, larger than or equal to the maximal fraction of data explained by model Δ. The natural reaction to this is to consider a penalization of model Δ that is monotonically dependent on the size of the model.[28]

Notice that the set of all stochastic choice functions can be obtained by taking the product of $|\mathcal{D}|$ simplices. In other words, the set of all stochastic choice functions can be seen as a subset of $[0, 1]^{|\mathcal{O}| - |\mathcal{D}|}$. Because all relevant stochastic models have a strictly lower dimensionality, they all have zero Lebesgue measure in the subspace of all stochastic choice functions. Therefore, any measure based on the Lebesgue volume of these models would regard all models as having the same size, and would differentiate them only in terms of the fraction of the data they rationalize.[29]

An alternative approach, in the spirit of the Akaike information criterion, would be to consider a cost dependent on the largest value of n, such that the model Δ has nonzero measure in the space $[0, 1]^n$. The essence of this perspective is to count the number of parameters in model Δ. Of the models previously analyzed, the deterministic choice model matches a finite subset of possible data sets and, hence, does not have a strictly positive dimension. The tremble model involves one tremble parameter and, hence, has dimension 1. The Luce model involves one utility value for each alternative and, when normalizing the sum of utility values, involves as many parameters as the number of alternatives minus 1. The SCRUM involves a probability measure over a subset of T preferences. If all menus are available, the dimension of this model is $T - 1$.

7.2. Convex Models

We have assumed model Δ to be closed, a basic property that guarantees the existence of maximal separations. An obvious further property to be considered is convexity, especially in relation to mixture models. These are common when dealing with heterogeneity at the population level, and can also be used to discuss intrapersonal heterogeneity. In a mixture model, the researcher convexifies a set of instances of a base model, allowing different subpopulations to be explained by different instances of the model. Notice that our methodology directly enables this type of analysis, because one can simply consider the desired, convexified, Δ model as the object of analysis.

---

28.   Another approach that would entail comparing the completeness of the different models, that is, the amount of predictive variation rationalized by the model. See Fudenberg, Kleinberg, Liang, and Mullainathan (2019) for a recent formal treatment of the notion of completeness.

29.   Another normalization that would not discriminate beyond absolute performance is $(\lambda^* - \lambda^{\min})/(\lambda^{\max} - \lambda^{\min})$, where $\lambda^{\max}$ and $\lambda^{\min}$ are a model's maximum and minimum performance values when studying all possible data sets. Clearly, $\lambda^{\max} = 1$ for all the models, and it can be easily shown that $\lambda^{\min} = 0$ for all the models discussed in this paper.

As an example of this approach, see the analysis of the SCRUM in Section 4.4, which can be understood as the convex hull of a subset of deterministic model instances.

The convexity of Δ may have useful implications. Given data ρ and model Δ, consider two separations $\langle \lambda, \delta, \varepsilon \rangle$ and $\langle \lambda', \delta', \varepsilon' \rangle$, and let $\alpha \in [0, 1]$. Clearly, $\alpha\lambda + (1 - \alpha)\lambda' \in [0, 1]$ and $\alpha\varepsilon + (1 - \alpha)\varepsilon' \in \mathrm{SCF}$, due to the convexity of $[0,1]$ and SCF. Whenever model Δ is convex, we also obtain that $\alpha\delta + (1 - \alpha)\delta' \in \Delta$, and hence, $\alpha\langle \lambda, \delta, \varepsilon \rangle + (1 - \alpha)\langle \lambda', \delta', \varepsilon' \rangle$ is also a separation, showing the convexity of the set of all separations. This transforms the search for maximal separations into a convex optimization problem.

It is important to note, however, that convex choice models are the exception rather than the norm. It is immediately obvious, for example, that the deterministic model is not convex. A mixture of two deterministic choice functions rationalized by two different preferences will clearly lead to a stochastic choice function that cannot be rationalized by any other preference. In a similar vein, it is well known that the Luce model represents another case of a nonconvex model (see Gul, Natenzon, and Pesendorfer 2014). Hence, the assumption of convexity, although not required for our results, would come with some loss of generality.

7.3. Residual Behavior

In our approach to finding the maximal fraction of the data consistent with a model, we have given the best possible chance to the model by leaving the space of possible residual behaviors completely unstructured. That is, we have assumed that residual behavior ε can be selected from the whole set of stochastic choice functions, SCF. Consequently, as the proof of Proposition 1 shows, a necessary condition for a separation to be maximal is that residual behavior lies exactly on the frontier of SCF. In other words, the residual behavior in a maximal separation imposes zero choice probabilities for some observations, which we call critical observations.

Sometimes the interest lies in separations involving less extreme residual behaviors, which might lead us to consider the possibility of imposing on the space of allowable residual behaviors a particular minimal structure beyond that of a stochastic choice function. The aim might be to consider the case in which residual behavior is in some way similar in nature to the reference model Δ, while allowing for more flexibility.

A set of minimal assumptions is sufficient to guarantee that the logic behind our methodology applies when considering restricted spaces of residual behavior, $\mathrm{RB} \subseteq \mathrm{SCF}$. In particular, we only need to consider that (i) the space of residual behaviors is a relaxation of the model, that is, $\Delta \subseteq \mathrm{RB}$; (ii) the data belong to the space of residual behaviors, that is, $\rho \in \mathrm{RB}$; and (iii) the space of residual behaviors has some technical properties, such as closedness and convexity, similar to those of the SCF space. Under these conditions, the concept of separation can be reformulated, provided that $\langle \lambda, \delta, \varepsilon \rangle \in [0, 1] \times \Delta \times \mathrm{RB}$. The logic of Proposition 1 remains valid, and a necessary condition for a separation to be maximal will be that residual behavior lies on the frontier of RB.

We conclude with some final comments on the interpretation of residual behavior, where we distinguish three cases. First, consider the situation in which the residual has

the consistency properties typical of a noisy structure. To illustrate, consider that the data ρ are generated exactly by the tremble model using a preference P and tremble γ, but the analyst initially approaches the data from the deterministic model perspective. The maximal separation will identify the true preference P and the residual will have a very transparent structure: The optimal alternative in menu A according to P is chosen with zero probability, whereas any other alternative is chosen with probability $1/(|A| - 1)$. Clearly, the structure of ε is very informative about the existing behavioral noise, and the analyst may wish to adopt the tremble model instead of the deterministic one.

Secondly, suppose that the residual has consistency properties typical of a competing instance of the model or of a competing model. To illustrate, consider that the data ρ are generated by a mixture of preferences P and P' (the former in larger proportion), but the analyst initially approaches the data from the perspective of the deterministic model. The maximal separation will identify preference P and the residual will have a very clear structure, that of preference P'. Clearly, the structure of ε is very informative about the existing heterogeneity, and, again, the analyst may wish to reconsider the choice of model for incorporating this heterogeneity into a mixture model. Similar reasoning applies when the residual resembles an instance not of model Δ but of some other reasonable model Δ'.

Finally, suppose that the residual is found to appear rather inconsistent. Here, a potentially fruitful option would be to apply the maximal separation approach on ε using some reasonable choice model, to assess the possibility of making any sense out of the apparently chaotic behavior ε. That is, try to ascertain whether ε itself can be understood, to a significant extent, as the combination of some choice model and another expression of residual behavior.

Appendix A: Proofs

Proof of Proposition 1

First, consider the case where $\rho \in \Delta$. Then, $\langle 1, \rho, \rho \rangle$ is clearly a maximal separation. Moreover, given that $\min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta(a, A)) = 1$ if and only if $\rho = \delta$, the result follows.

Let us now consider the case of $\rho \notin \Delta$. We start by claiming that, for a given $\delta \in \Delta$, there exist $\lambda \in [0, 1)$ and $\varepsilon \in \text{SCF}$ such that $\langle \lambda, \delta, \varepsilon \rangle$ is a separation if and only if $\lambda \leq \min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta(a, A))$. To prove the "only if" part, assume that $\langle \lambda, \delta, \varepsilon \rangle$ is a separation. Then, it must be the case that $\rho = \lambda\delta + (1 - \lambda)\varepsilon$, or, equivalently, $(\rho - \lambda\delta)/(1 - \lambda) = \varepsilon \geq 0$. This implies that $\rho - \lambda\delta \geq 0$ and, ultimately, that $\lambda \leq \rho/\delta$. Hence, it must be that $\lambda \leq \min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta(a, A))$, as desired.[30] To prove the "if" part, suppose that $\lambda \leq \min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta(a, A))$. We

---

30. Notice that, in dividing by δ, we are using the above-mentioned convention.

now prove that $\langle \lambda, \delta, \varepsilon = (\rho - \lambda\delta)/(1 - \lambda) \rangle$ is a separation of the data. Because, by assumption, $\delta \in \Delta$ and the construction guarantees that $\rho = \lambda\delta + (1 - \lambda)\varepsilon$, we are only required to prove that $\varepsilon \in \text{SCF}$. We begin by checking that $\varepsilon(a, A) \geq 0$ holds for every $(a, A) \in \mathcal{O}$. To see this, suppose by contradiction that this is not true. Then, there would exist $(b, B) \in \mathcal{O}$ such that $(\rho(b, B) - \lambda\delta(b, B))/(1 - \lambda) < 0$. This would imply that $\rho(b, B) - \lambda\delta(b, B) < 0$ and, hence, that $\delta(b, B) > 0$, with $\rho(b, B)/\delta(b, B) < \lambda \leq \min\limits_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta(a, A))$, which is a contradiction. Finally, it is also the case that $\sum_{a\in A} \varepsilon(a, A) = \sum_{a\in A}(\rho(a, A) - \lambda\delta(a, A))/(1 - \lambda) = (1 - \lambda)/(1 - \lambda) = 1$ for every $A \in \mathcal{D}$. Therefore, $\varepsilon \in \text{SCF}$ and the claim is proved.

Now, the aforementioned claim shows that the maximal fraction that can be explained with model $\{\delta\}$ is $\min\limits_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta(a, A))$. This argument immediately implies the desired results on Δ, provided that maximal separations exist.

We now show the existence of maximal separations. Given the domain, any separation $\langle \lambda, \delta, \varepsilon \rangle$ of ρ is a vector in $\mathbb{R}^n$, with $n = 2|\mathcal{O}| + 1$. We first prove that the set of separations is a closed subset of $\mathbb{R}^n$. Consider a sequence of separations $\langle \lambda_t, \delta_t, \varepsilon_t \rangle_{t=1}^{\infty}$ and suppose that this sequence converges in $\mathbb{R}^n$. Given the finite dimensionality and the closure of Δ and SCF, we clearly have that $\lim_t \lambda_t \in [0, 1]$, $\lim_t \delta_t \in \Delta$, and $\lim_t \varepsilon_t \in \text{SCF}$, and it is evident that $\langle \lim_t \lambda_t, \lim_t \delta_t, \lim_t \varepsilon_t \rangle$ is a separation of ρ. This proves that the set of separations is closed, and, being a subset of $[0, 1]^n$, it is also bounded and, hence, compact. Because the maximal fraction of data explained can be thought of as the result of maximizing, over the set of separations, the projection map assigning the first component of the separation, that is, value λ, existence is guaranteed. □

Proof of Proposition 2

Let $\{\lambda_S\}_{S:\mathcal{D}|_S\neq\varnothing}$ and $P \in \mathcal{P}$ satisfy (1) and (2). For every S such that $\mathcal{D}|_S \neq \varnothing$, denote by $\text{DET}_{\mathcal{D}|_S}$ the deterministic rational stochastic choice functions defined over the subdomain $\mathcal{D}|_S$. Similarly, denote by $\rho|_S$ the restriction of ρ to $\mathcal{D}|_S$. We start by proving, recursively, that the maximal fraction of data $\rho|_S$ explained by model $\text{DET}_{\mathcal{D}|_S}$ is equal to λ_S. Consider any subset S for which $\mathcal{D}|_S = \{S\}$. In this case, Proposition 1 guarantees that the maximal fraction of data $\rho|_S$ explained by model $\text{DET}_{\mathcal{D}|_S}$ is

$$\max_{\delta\in\text{DET}_{\mathcal{D}|_S}} \min_{(a,A)\in\mathcal{O}|_S} \frac{\rho|_S(a, A)}{\delta(a, A)} = \max_{P\in\mathcal{P}} \min_{(a,A)\in\mathcal{O}|_S} \frac{\rho(a, A)}{\delta_P(a, A)} = \max_{P\in\mathcal{P}} \min_{a\in S} \frac{\rho(a, S)}{\delta_P(a, S)}$$

$$= \max_{P\in\mathcal{P}} \frac{\rho(m_P(S), S)}{\delta_P(m_P(S), S)} = \max_{P\in\mathcal{P}} \rho(m_P(S), S)$$

$$= \max_{a\in S} \rho(a, S) = \max_{a\in S} \min_{(a,A)\in\mathcal{O}|_S} \rho(a, A) = \lambda_S.$$

Now suppose that $\mathcal{D}|_S \neq \{S\}$ and that the result has been proven for any strict subset of S with a nonempty subdomain. For any $a \in S$, denote by $\mathcal{P}_{aS}$ the set of preferences that rank a above any other alternative in S, that is,

$\mathcal{P}_{aS} = \{P \in \mathcal{P} : a = m_P(S)\}$, and by $\mathtt{aS}$ the subset of $\mathtt{DET}_{\mathcal{D}|_S}$ generated by preferences in $\mathcal{P}_{aS}$. Trivially, $\mathtt{DET}_{\mathcal{D}|_S} = \cup_{a \in S} \mathtt{aS} = \cup_{a \in S} \cup_{P \in \mathcal{P}_{aS}} \{\delta_P\}$. Because the only observations for which δ_P has a nonnull value are those that take form $(m_P(A), A)$, Proposition 1 guarantees that the maximal fraction of data $\rho|_S$ explained by model $\mathtt{DET}_{\mathcal{D}|_S}$ is $\max_{a \in S} \max_{P \in \mathcal{P}_{aS}} \min_{A \in \mathcal{D}|_S} \rho(m_P(A), A)$. Because $P \in \mathcal{P}_{aS}$, we obtain that $m_P(A) = a$ whenever $a \in A$, and, hence, the latter value is equal to $\max_{a \in S} \max_{P \in \mathcal{P}_{aS}} \min\{\{\rho(a, A)\}_{(a,A) \in \mathcal{O}|_S}, \{\rho(m_P(B), B)\}_{B \in \mathcal{D}|_{S \setminus \{a\}}}\}$. This can be expressed as $\max_{a \in S} \min\{\{\rho(a, A)\}_{(a,A) \in \mathcal{O}|_S}, \max_{P \in \mathcal{P}_{aS}} \min_{B \in \mathcal{D}|_{S \setminus \{a\}}} \rho(m_P(B), B)\}$ or, equivalently, as $\max_{a \in S} \min\{\{\rho(a, A)\}_{(a,A) \in \mathcal{O}|_S}, \min_{B \in \mathcal{D}|_{S \setminus \{a\}}} \max_{P \in \mathcal{P}_{aS}} \min_{C \in \mathcal{D}|_B} \rho(m_P(C), C)\}$. Given that $a \notin B$, it is clearly the case that $\max_{P \in \mathcal{P}_{aS}} \min_{C \in \mathcal{D}|_B} \rho(m_P(C), C) = \max_{P \in \mathcal{P}} \min_{C \in \mathcal{D}|_B} \rho(m_P(C), C)$, and, by Proposition 1 and the structure of deterministic stochastic choice functions, the latter is the maximal fraction of data $\rho|_B$ explained by model $\mathtt{DET}_{\mathcal{D}|_B}$, which is equal to λ_B by hypothesis. Hence, the maximal fraction of data $\rho|_S$ explained by model $\mathtt{DET}_{\mathcal{D}|_S}$ must also be equal to λ_S, as desired. As a corollary, we have that the maximal fraction of the data explained by the deterministic model is λ_X and the claim follows from the construction. □

Proof of Proposition 3

Because the proof has the same structure as the proof of Proposition 2, we skip some of the steps and use the same notation as before. We start by (recursively) proving that the maximal fraction of data $\rho|_S$ explained by the collection of stochastic choice functions in $\mathtt{Tremble}_{\mathcal{D}|_S}$ with a fixed degree of tremble γ, which we denote by $\mathtt{Tremble}_{\mathcal{D}|_S}(\gamma)$, is equal to $\lambda_S(\gamma)$. We start with any subset S for which $\mathcal{D}|_S = \{S\}$. The maximal fraction of data $\rho|_S$ explained by $\mathtt{Tremble}_{\mathcal{D}|_S}(\gamma)$ is

$$\max_{\delta \in \mathtt{Tremble}_{\mathcal{D}|_S}(\gamma)} \min_{(a,A) \in \mathcal{O}|_S} \frac{\rho|_S(a, A)}{\delta(a, A)}$$

$$= \max_{P \in \mathcal{P}} \min\left\{\frac{\rho(m_P(S), S)}{\delta_{[P,\gamma]}(m_P(S), S)}, \left\{\frac{\rho(b, S)}{\delta_{[P,\gamma]}(b, S)}\right\}_{b \in S \setminus \{m_P(S)\}}\right\}$$

$$= \max_{P \in \mathcal{P}} \min\left\{\frac{|S|\rho(m_P(S), S)}{(1-\gamma)|S| + \gamma}, \left\{\frac{|S|\rho(b, S)}{\gamma}\right\}_{b \in S \setminus \{m_P(S)\}}\right\}$$

$$= \max_{a \in S} \min\left\{\frac{|S|\rho(a, S)}{(1-\gamma)|S| + \gamma}, \left\{\frac{|S|\rho(b, S)}{\gamma}\right\}_{b \in S \setminus \{a\}}\right\}$$

$$= \max_{a \in S} \min\left\{\left\{\frac{|A|\rho(a, A)}{(1-\gamma)|A| + \gamma}\right\}_{(a,A) \in \mathcal{O}|_S}, \left\{\frac{|A|\rho(b, A)}{\gamma}\right\}_{\substack{(b,A) \in \mathcal{O}|_S \\ b \neq a}}\right\}$$

$$= \lambda_S(\gamma).$$

Whenever $\mathcal{D}|_S \neq \{S\}$, we can write the maximal fraction of data $\rho|_S$ explained by model $\mathtt{Tremble}_{\mathcal{D}|_S}(\gamma)$ as $\max_{a \in S} \max_{P \in \mathcal{P}_{aS}} \min \{\{|A|\rho(m_P(A), A)/((1 - \gamma)|A| + \gamma)\}_{A \in \mathcal{D}|_S}, \{|A|\rho(b, A)/\gamma\}_{(b,A) \in \mathcal{O}|_S, b \neq m_P(A)}\}$. Notice that we can decompose $\{|A|\rho(m_P(A), A)/((1 - \gamma)|A| + \gamma)\}_{A \in \mathcal{D}|_S}$ into $\{|A|\rho(a, A)/((1 - \gamma)|A| + \gamma)\}_{(a,A) \in \mathcal{O}|_S}$ and $\{|B|\rho(m_P(B), B)/((1 - \gamma)|B| + \gamma)\}_{B \in \mathcal{D}|_{S \setminus \{a\}}}$. Similarly, we can decompose $\{|A|\rho(b, A)/\gamma\}_{(b,A) \in \mathcal{O}|_S, b \neq m_P(A)}$ into components $\{|A|\rho(b, A)/\gamma\}_{(b,A) \in \mathcal{O}|_S, b \neq a \in A}$ and $\{|B|\rho(b, B)/\gamma\}_{B \in \mathcal{D}|_{S \setminus \{a\}}, b \neq m_P(B)}$. By the same reasoning as in the proof of Proposition 2, consideration of both $\{|B|\rho(m_P(B), B)/((1 - \gamma)|B| + \gamma)\}_{B \in \mathcal{D}|_{S \setminus \{a\}}}$ and $\{|B|\rho(b, B)/\gamma\}_{B \in \mathcal{D}|_{S \setminus \{a\}}, b \neq m_P(B)}$ yields the value $\{\lambda_B(\gamma)\}_{B \in \mathcal{D}|_{S \setminus \{a\}}}$. This proves the claim. From Proposition 2, the maximal separations for model $\mathtt{Tremble}_{\mathcal{D}}(\gamma)$ explain a fraction $\lambda_X(\gamma)$ of the data. Because $\mathtt{Tremble}_{\mathcal{D}} = \cup_\gamma \mathtt{Tremble}_{\mathcal{D}}(\gamma)$, one simply needs to consider the value γ^* maximizing $\lambda_X(\gamma)$, and the result follows immediately from Proposition 1. $\qquad\square$

Proof of Proposition 4

To prove the "if" part, let $\delta_L \in \mathtt{Luce}$ and suppose that there exists $\{(a_i, A_i)\}_{i=1}^I \subseteq \mathcal{O}_{\delta_L}$ such that $\bigcup_{i=1}^I \{a_i\} = \bigcup_{i=1}^I A_i$. From Proposition 1, the maximal fraction that can be explained by model $\{\delta_L\}$ is $\min_{(a,A) \in \mathcal{O}} (\rho(a, A)/\delta_L(a, A))$. Assume, by way of contradiction, that δ_L is not part of a maximal separation for the Luce model. Therefore, there exists $\langle \lambda^*, \delta_L^*, \varepsilon^* \rangle$ such that, for every $i \in \{1, 2, \ldots, I\}$, $\rho(a_i, A_i)/\delta_L(a_i, A_i) = \min_{(a,A) \in \mathcal{O}} (\rho(a, A)/\delta_L(a, A)) < \lambda^* = \min_{(a,A) \in \mathcal{O}} (\rho(a, A)/\delta_L^*(a, A)) \leq \rho(a_i, A_i)/\delta_L^*(a_i, A_i)$. For every $i \in \{1, 2, \ldots, I\}$, we have that $\rho(a_i, A_i) > 0$, and, hence, because the ρ/δ_L ratio is minimized at $\mathcal{O}_{\delta_L}$, it must be that $\delta_L(a_i, A_i) > 0$, making $\rho(a_i, A_i)/\delta_L(a_i, A_i) < \rho(a_i, A_i)/\delta_L^*(a_i, A_i)$ equivalent to $\delta_L^*(a_i, A_i) < \delta_L(a_i, A_i)$. Let $\{\delta_{v_n}'\}_{n=1}^\infty$ and $\{\delta_{u_n}\}_{n=1}^\infty$ be two sequences of strictly positive Luce stochastic choice functions that converge to δ_L^* and δ_L, respectively. Select an m sufficiently large that $\delta_L^*(a_i, A_i) < \delta_{u_m}(a_i, A_i)$ holds for every $i \in \{1, 2, \ldots, I\}$. Given m, now select an m' sufficiently large that, for every $i \in \{1, 2, \ldots, I\}$, $\delta_{v_{m'}}'(a_i, A_i) < \delta_{u_m}(a_i, A_i)$ holds. We then have that $1/(\sum_{x \in A_i} v_{m'}(x)/v_{m'}(a_i)) = v_{m'}(a_i)/(\sum_{x \in A_i} v_{m'}(x)) = \delta_{v_{m'}}'(a_i, A_i) < \delta_{u_m}(a_i, A_i) = u_m(a_i)/(\sum_{x \in A_i} u_m(x)) = 1/(\sum_{x \in A_i} u_m(x)/u_m(a_i))$, thus guaranteeing, for every $i \in \{1, 2, \ldots, I\}$, the existence of one alternative $\bar{x}_i \in A_i \setminus \{a_i\}$ such that $v_{m'}(a_i)/v_{m'}(\bar{x}_i) < u_m(a_i)/u_m(\bar{x}_i)$. Given that $\bigcup_{i=1}^I \{a_i\} = \bigcup_{i=1}^I A_i$, there exists a subcollection $\{a_{i_h}\}_{h=1}^H$ of $\{a_i\}_{i=1}^I$ with the following properties: (i) $a_{i_{h+1}} \in A_{i_h}$, with $h = 1, \ldots, H - 1$, and $a_{i_1} \in A_{i_H}$; and (ii) $v_{m'}(a_{i_h})/v_{m'}(a_{i_{h+1}}) < u_m(a_{i_h})/u_m(a_{i_{h+1}})$, with $h = 1, \ldots, H - 1$, and

$v_{m'}(a_{i_H})/v_{m'}(a_{i_1}) < u_m(a_{i_H})/u_m(a_{i_1})$. Obviously,

$$1 = \frac{v_{m'}(a_{i_H})}{v_{m'}(a_{i_1})} \Pi_{h=1}^{H-1} \frac{v_{m'}(a_{i_h})}{v_{m'}(a_{i_{h+1}})} < \frac{u_m(a_{i_H})}{u_m(a_{i_1})} \Pi_{h=1}^{H-1} \frac{u_m(a_{i_h})}{u_m(a_{i_{h+1}})} = 1,$$

which is a contradiction. This concludes the "if" part of the proof.

To prove the "only if" part, suppose that δ_L belongs to a maximal separation for the Luce model. Let $[x]$ be the set of all alternatives $x' \in X$ for which there exists a sequence of observations $\{(b_j, B_j)\}_{j=1}^J$, with (i) $x = b_1$ and $x' \equiv b_{J+1} \in B_J$; and (ii) for every $j \in \{1, 2, \ldots, J\}$, $\delta_L(b_j, B_j) > 0$ and $\delta_L(b_{j+1}, B_j) > 0$. If there is no alternative for which such a sequence exists, let $[x] = \{x\}$. Clearly, $[\cdot]$ defines equivalence classes on X. Whenever there exists $A \in \mathcal{D}$ with $\{x, y\} \subseteq A$ and $\delta_L(x, A) > \delta_L(y, A) = 0$, we write $[x] \succ [y]$. We claim that $\succ$ is an acyclic relation on the set of equivalence classes. To see this, assume, by contradiction, that there is a cycle of pairs $\{a_q, b_q\}$, menus $A_q \supseteq \{a_q, b_q\}$, and equivalence classes $[x_q]$, $q \in \{1, 2, \ldots, Q\}$, such that (i) $\delta_L(a_q, A_q) > \delta_L(b_q, A_q) = 0$ for every $q \in \{1, 2, \ldots, Q\}$, (ii) $a_q \in [x_q]$ for every $q \in \{1, 2, \ldots, Q\}$; and (iii) $b_q \in [x_{q+1}]$ for every $q \in \{1, 2, \ldots, Q-1\}$ and $b_Q \in [x_1]$. We can then consider a sequence of stochastic choice functions $\{\delta_{u_n}\}_{n=1}^\infty$ that converges to δ_L. Because b_q and a_{q+1} belong to the same equivalence class $[x_{q+1}]$, either $b_q = a_{q+1}$ or there exists a sequence of observations $\{(d_j, D_j)\}_{j=1}^J$ with (i) $b_q = d_1$ and $a_{q+1} = d_{J+1} \in D_J$; and (ii) for every $j \in \{1, 2, \ldots, J\}$, $\delta_L(d_j, D_j) > 0$ and $\delta_L(d_{j+1}, D_j) > 0$ (and the same holds for a_Q and b_1). Define the strictly positive constant $K_q = 1$ whenever $b_q = a_{q+1}$, and $K_q = (1/2)\Pi_{j=1}^J(\delta_L(d_j, D_j)/\delta_L(d_{j+1}, D_j))$ otherwise (with a similar definition for K_Q relating a_Q and b_1). If $b_q = a_{q+1}$, then, trivially, $u_n(b_q) = u_n(a_{q+1})$ for every n. Otherwise, for a sufficiently large n in the sequence $\{u_n\}_{n=1}^\infty$, we have that $u_n(b_q)/u_n(a_{q+1}) = \Pi_{j=1}^J(u_n(d_j)/u_n(d_{j+1})) = \Pi_{j=1}^J(\delta_{u_m}(d_j, D_j)/\delta_{u_m}(d_{j+1}, D_j)) \geq K_q$. Hence, in any case, $u_n(b_q)/K_q \geq u_n(a_{q+1})$ holds for any sufficiently large n (and the same holds for b_Q and a_1). Also, because $\delta_L(a_q, A_q) > \delta_L(b_q, A_q) = 0$ for every $q \in \{1, 2, \ldots, Q\}$, we can find an n sufficiently large that $u_n(a_q) > u_n(b_q)/K_q$. Hence, we can find an m that is sufficiently large that $u_m(a_1) > u_m(b_1)/K_1 \geq u_m(a_2) > u_m(b_2)/K_2 \geq \ldots \geq u_m(a_Q) > u_m(b_Q)/K_Q \geq u_m(a_1)$. This is a contradiction, which proves the acyclicity of $\succ$. We can then denote the equivalence classes as $\{[x_e]\}_{e=1}^E$, where $[x_e] \succ [x_{e'}]$ implies that $e < e'$. For an equivalence class $[x_e]$, define the vector $u_{[x_e]} \in \mathcal{U}$ such that $u_{[x_e]}(y) = 0$ if $y \notin [x_e]$ and, $u_{[x_e]}(y)/u_{[x_e]}(y') = \delta_L(y, A)/\delta_L(y', A)$ whenever $y, y' \in [x_e]$, $\delta_L(y, A) > 0$ and $\delta_L(y', A) > 0$. This is clearly well defined due to the structure of Luce stochastic choice functions. Now consider the sequence of Luce stochastic choice functions $\{\delta_{v_n}\}_{n=1}^\infty$ given by $v_n = (1 - \sum_{e=2}^E(1/2^e)^n)u_{[x_1]} + \sum_{e=2}^E(1/2^e)^n u_{[x_e]}$, which clearly converges to δ_L. Consider the following three collections of observations: $\mathcal{O}_1$, $\mathcal{O}_2$, and $\mathcal{O}_3$. $\mathcal{O}_1$ is composed of all observations $(a, A) \in \mathcal{O}$ such that $A \subseteq [a]$. $\mathcal{O}_2$ is composed of all observations $(a, A) \in \mathcal{O} \setminus \mathcal{O}_1$, such that $b \in A$, $a \in [a_i]$ and $b \in [a_j]$ imply $i \geq j$. $\mathcal{O}_3$ is composed of observations in $\mathcal{O} \setminus (\mathcal{O}_1 \cup \mathcal{O}_2)$. Notice that, for an n sufficiently large, for every $(a, A) \in \mathcal{O}_1$, we have that

$\rho(a, A)/\delta_{v_n}(a, A) = \rho(a, A)/\delta_L(a, A)$, and for every $(a, A) \in \mathcal{O}_2$, we have that $\rho(a, A)/\delta_{v_n}(a, A) > \rho(a, A)/\delta_L(a, A)$. Also, for an n sufficiently large, $(1/2)^n < \min\{\rho(a, A) : a \in A \in \mathcal{D}\}$, and, hence, $(a, A) \in \mathcal{O}_3$ implies that $\rho(a, A)/\delta_{v_n}(a, A) \geq \rho(a, A)/((1/2)^m) > 1$. In this case, we can fix an m sufficiently large that, from Proposition 1, $\min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta_{v_m}(a, A)) = \min_{(a,A)\in\mathcal{O}_1\cup\mathcal{O}_2} (\rho(a, A)/\delta_{v_m}(a, A)) \geq \min_{(a,A)\in\mathcal{O}_1\cup\mathcal{O}_2} (\rho(a, A)/\delta_L(a, A)) \geq \min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta_L(a, A))$.[31] Indeed, because δ_L belongs to a maximal separation of the Luce model, it must be that $\min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta_{v_m}(a, A)) = \min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta_L(a, A))$, and, hence, $\mathcal{O}_1$ is nonempty, with $\mathcal{O}_{\delta_{v_m}} \subseteq \mathcal{O}_{\delta_L} \subseteq \mathcal{O}_1$.

Assume, by way of contradiction, that there is no subcollection $\{(a_i, A_i)\}_{i=1}^I \subseteq \mathcal{O}_{\delta_L}$ such that $\bigcup_{i=1}^I \{a_i\} = \bigcup_{i=1}^I A_i$. Then, for every subcollection $\{(a_i, A_i)\}_{i=1}^I \subseteq \mathcal{O}_{\delta_{v_m}}$ it must also be that $\bigcup_{i=1}^I \{a_i\} \neq \bigcup_{i=1}^I A_i$. Hence, there must exist at least one alternative x such that $x \neq a$ for every $(a, A) \in \mathcal{O}_{\delta_{v_m}}$ and $x \in A$ for some $(a, A) \in \mathcal{O}_{\delta_{v_m}}$. Consider the segment $\alpha 1_x + (1 - \alpha)v_m$, with $\alpha \in [0, 1]$. Select the maximal separation in this segment, which can be identified as follows. Partition the set of observations into two classes $\mathcal{O}' = \{(a, A) \in \mathcal{O}, a \neq x \in A\}$ and $\mathcal{O}'' = \mathcal{O} \setminus \mathcal{O}'$ and then select the Luce utilities defined by the unique value $\bar{\alpha} \in [0, 1]$ that solve $\min_{(a,A)\in\mathcal{O}'} (\rho(a, A)/\delta_{\alpha 1_x + (1-\alpha)v_m}(a, A)) = \min_{(a,A)\in\mathcal{O}''} (\rho(a, A)/\delta_{\alpha 1_x + (1-\alpha)v_m}(a, A))$. Notice that, given the structure of the Luce model, the left-hand ratio increases with α, continuously and strictly, approaching infinity. At the same time, the right-hand ratio weakly decreases with α continuously. Notice also that, for $\alpha = 0$, the left-hand ratio is strictly lower than the right-hand ratio. This is because there exists at least one observation on the left-hand side that belongs to $\mathcal{O}_{\delta_{v_m}}$. Thus, $\bar{\alpha}$ must exist, and Proposition 1 guarantees that this provides the maximal separation in the segment. Then, consider the vector of Luce utilities $v = \bar{\alpha} 1_x + (1 - \bar{\alpha})v_m$. If alternative x is present in all the menus in $\mathcal{O}_{\delta_{v_m}}$, then $\min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta_v(a, A)) > \min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta_{v_m}(a, A)) = \min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta_L(a, A))$, thus contradicting the maximality of δ_L. If x is not present in some menu of $\mathcal{O}_{\delta_{v_m}}$, it must be the case that $\mathcal{O}_{\delta_v} \subsetneq \mathcal{O}_{\delta_{v_m}}$ and $\min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta_v(a, A)) = \min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta_L(a, A))$. Given the finiteness of the data, we can repeat the same exercise for δ_v and, eventually, contradict the optimality of δ_L. This concludes the proof.   □

---

31.   This further shows that there is always a strictly positive instance of Luce that is maximal.

Proof of Proposition 5

We start by proving that λ_T is lower than or equal to the maximal fraction of the data that can be explained by SCRUM. By construction, it is guaranteed that $1 \geq \lambda_T \geq \lambda_{T-1} \geq \ldots \geq \lambda_1 \geq 0$. Whenever $\lambda_T = 0$, the result is immediate. Assume that $\lambda_T \in (0, 1)$. We prove that there exists a separation of ρ of the form $\langle \lambda_T, \delta_{\mu_T}, (\rho - \lambda_T \delta_{\mu_T})/(1 - \lambda_T) \rangle$. Because the construction guarantees that $\delta_{\mu_T} \in$ SC, we only need to prove that $\varepsilon = (\rho - \lambda_T \delta_{\mu_T})/(1 - \lambda_T) \in$ SCF. To see this, consider $(a, A) \in \mathcal{O}$ and denote by $\underline{i}$ and $\overline{i}$ the integers of the first and last preferences in $\mathcal{P}'$, such that a is the maximal element in A. The construction also guarantees that $\rho(a, A) \geq \lambda_{\overline{i}} - \lambda_{\underline{i}-1} = \lambda_T(\lambda_{\overline{i}} - \lambda_{\underline{i}-1})/\lambda_T$. Now, the recursive equations can be written as $\mu_T(P_i) = (\lambda_{P_i} - \lambda_{P_{i-1}})/\lambda_T$ for every $i \in \{1, 2, \ldots, T\}$, with $\lambda_0 = 0$ and, hence, $\rho(a, A) \geq \lambda_T \sum_{i=\underline{i}}^{\overline{i}} \mu_T(P_i) = \lambda_T \delta_{\mu_T}(a, A)$. This implies that $\varepsilon(a, A) \geq 0$. Notice also that $\sum_{a \in A} \varepsilon(a, A) = \sum_{a \in A}(\rho(a, A) - \lambda_T \delta_{\mu_T}(a, A))/(1 - \lambda_T) = (1 - \lambda_T)/(1 - \lambda_T) = 1$, thus proving that $\varepsilon \in$ SCF. This shows the claim and, hence, the desired inequality. Finally, suppose that $\lambda_T = 1$. In this case, note, again, that the construction guarantees that $\rho = \delta_{\mu_T} \in$ SC, and the desired inequality follows.

We now show that λ_T is greater than or equal to the maximal fraction of the data that can be explained by SCRUM. To show this, let $\langle \lambda, \delta_\mu, \varepsilon \rangle$ be a separation for SCRUM. We need to show that $\lambda_T \geq \lambda$. We proceed recursively to show that $\lambda_i \geq \sum_{j=1}^{i} \lambda \mu(P_j)$ holds, and, hence, $\lambda_T \geq \sum_{j=1}^{T} \lambda \mu(P_j) = \lambda$, as desired. Let $i = 1$ and A' be a menu solving $\min_{A \in \mathcal{D}} \rho(m_{P_1}(A), A)$. Hence, $\lambda_1 - \lambda \mu(P_1) = \rho(m_{P_1}(A'), A') - \lambda \mu(P_1) \geq \rho(m_{P_1}(A'), A') - \lambda \sum_{j:m_{P_j}(A')=m_{P_1}(A')} \mu(P_j)$. By the definition of SCRUMs, the last expression can be written as $\rho(m_{P_1}(A'), A') - \lambda \delta_\mu(m_{P_1}(A'), A')$, or, equivalently, as $(1 - \lambda)\varepsilon(m_{P_1}(A'), A')$. Because $\varepsilon \in$ SCF, the latter expression must be positive, thus proving the desired result. Suppose that the inequality is true for every P_j with $j < i$. We now prove this for P_i. Let $\overline{A}$ be a menu solving $\min_{A \in \mathcal{D}}[\rho(m_{P_i}(A), A) + \max_{j:j \leq i, m_{P_j}(A) \neq m_{P_i}(A)} \lambda_j]$. Then, we have $\rho(m_{P_i}(\overline{A}), \overline{A}) = \lambda \delta_\mu(m_{P_i}(\overline{A}), \overline{A}) + (1 - \lambda)\varepsilon(m_{P_i}(\overline{A}), \overline{A}) \geq \lambda \delta_\mu(m_{P_i}(\overline{A}), \overline{A}) = \lambda \sum_{P:m_P(\overline{A})=m_{P_i}(\overline{A})} \mu(P)$. If it is the case that $\{P : m_P(\overline{A}) = m_{P_i}(\overline{A})\} \supseteq \{P_1, P_2, \ldots, P_i\}$, then clearly $\lambda_i = \rho(m_{P_i}(\overline{A}), \overline{A}) \geq \lambda \sum_{P:m_P(\overline{A})=m_{P_i}(\overline{A})} \mu(P) = \sum_{j=1}^{i} \lambda \mu(P_j)$, and we have concluded the induction argument. Otherwise, the single-crossing condition guarantees that there exists $\overline{j} \in \{1, \ldots, i-1\}$ such that $\{P : m_P(\overline{A}) = m_{P_i}(\overline{A})\} \supseteq \{P_{\overline{j}+1}, P_{\overline{j}+2}, \ldots, P_i\}$ and $\rho(m_{P_i}(\overline{A}), \overline{A}) \geq \sum_{j=\overline{j}+1}^{i} \lambda \mu(P_j)$. In this case, the induction hypothesis also guarantees that $\lambda_{\overline{j}} \geq \sum_{j=1}^{\overline{j}} \lambda \mu(P_j)$. By combining these two inequalities, we are able to conclude that $\lambda_i \geq \sum_{j=1}^{i} \lambda \mu(P_j)$, and the induction step is complete. This implies, in particular, that $\lambda \leq \lambda_T$.

By combining the aforementioned two claims, we have shown that $\langle \lambda_T, \delta_{\mu_T}, (\rho - \lambda_T \delta_{\mu_T})/(1 - \lambda_T) \rangle$ is a maximal separation for SCRUM, which concludes the proof. $\square$

TABLE B.1. A stochastic choice function ρ.

	x	y	z
$\{x, y, z\}$	0.15	0.6	0.25
$\{x, y\}$	0.25	0.75	
$\{x, z\}$	0.7		0.3
$\{y, z\}$		0.4	0.6

Appendix B: Examples

We first propose a simple example of a stochastic choice function, and derive the maximal separations for all the models studied in Section 4. We then use another example to show that the maximal separations for the deterministic model and the tremble model do not necessarily identify the same preference relations. Finally, we propose a particular data-generating process and use the tremble model to illustrate our conjecture on the differences between maximal separation and maximum likelihood in the overestimation of choice probabilities.

Table B.1 reports a stochastic choice function ρ defined on every nonsingleton subset of $X = \{x, y, z\}$, that is, $\mathcal{D} = \{\{x, y, z\}, \{x, y\}, \{x, z\}, \{y, z\}\}$. Note that this stochastic choice function involves behavior that is rather unstructured, in the sense that it does not satisfy weak stochastic transitivity.

We start with the deterministic model, where we can first calculate the maximal fraction for every set for which $\mathcal{D}|_S = \{S\}$, that is, the binary sets:

$$\lambda_{\{x,y\}} = \max\{\rho(x, \{x, y\}), \rho(y, \{x, y\})\} = 0.75;$$

$$\lambda_{\{x,z\}} = \max\{\rho(x, \{x, z\}), \rho(z, \{x, z\})\} = 0.7; \text{ and}$$

$$\lambda_{\{y,z\}} = \max\{\rho(y, \{y, z\}), \rho(z, \{y, z\})\} = 0.6.$$

We can then proceed to assign a value to menu X, for which we first analyze the alternatives in X one-by-one, computing a minimum value for each, as follows. For alternative x, $\left\{\{\rho(x, \{x, y\}), \rho(x, \{x, z\}), \rho(x, X)\}, \lambda_{\{y,z\}}\right\} = \rho(x, \{x, y, z\}) = 0.15$; for alternative y, $\left\{\{\rho(y, \{x, y\}), \rho(y, \{y, z\}), \rho(y, X)\}, \lambda_{\{x,z\}}\right\} = \rho(y, \{y, z\}) = 0.4$; and for alternative z, $\left\{\{\rho(z, \{x, z\}), \rho(z, \{y, z\}), \rho(z, X)\}, \lambda_{\{x,y\}}\right\} = \rho(z, \{x, y, z\}) = 0.25$. Thus, we get

$$\lambda_X = \max\{0.15, 0.4, 0.25\} = 0.4.$$

Notice that the final value is the same as that obtained with alternative y. In subset $X\backslash\{y\}$, the alternative determining the value $\lambda_{\{x, z\}}$ is x. Hence, the second part of Proposition 2 guarantees that δ_P with $yPxPz$ conforms to a maximal separation of ρ. From $\lambda_X = 0.4$, one can immediately obtain the corresponding residual behavior as $\varepsilon = (\rho - 0.4\delta_P)/0.6$, that is, $\varepsilon(x, X) = 1/4$, $\varepsilon(y, X) = 1/3$, $\varepsilon(x, \{x, y\}) = 5/12$, $\varepsilon(x, \{x, z\}) = 1/2$, and $\varepsilon(y, \{y, z\}) = 0$. To close the discussion of this example, notice from the residual behavior that the frontier of SCF is reached at $(y, \{y, z\})$. This is precisely

the observation where the identified instance δ_P fails most seriously. It also determines the maximal fraction of data explained by DET, that is, $\rho(y, \{y, z\})/\delta_P(y, \{y, z\}) = 0.4/1 = 0.4$.

We now illustrate the treatment of the tremble model. After replicating the steps taken in the analysis of DET, we conclude that $yPxPz$ is the optimal preference relation for every given value of γ.[32] In seeking the optimal value of γ, note that there are only two possible critical observations, depending on the value of γ. When γ is low, we know, from the study of the deterministic case, that the critical observation is $(y, \{y, z\})$, with a ratio of ρ to δ equal to $0.4/(1 - \gamma + \gamma/2)$. When γ is high, the critical observation is $(x, \{x, y, z\})$, with a ratio of ρ to δ equal to $0.15/(\gamma/3)$. By noticing that the first ratio is increasing and starts at a value less than the second ratio, which is decreasing, it follows that the maximal fraction of data explained by the optimal tremble can be found by equating these two ratios, which yields $\gamma^* = 0.72$. Hence, the maximal fraction of data explained is 0.625, obtained with the trembling stochastic choice function $\delta_{[P, 0.72]}$ and residual behavior $\varepsilon = (\rho - 0.625\delta_{[P, 0.72]})/0.375$, that is, $\varepsilon(x, X) = 0$, $\varepsilon(y, X) = 11/15$, $\varepsilon(x, \{x, y\}) = 1/15$, $\varepsilon(x, \{x, z\}) = 4/5$, and $\varepsilon(y, \{y, z\}) = 0$. To conclude, notice that the explanatory power of the tremble model is limited by the tension created by the two critical observations, $(y, \{y, z\})$ and $(x, \{x, y, z\})$.

As for the Luce model, consider the Luce utilities $u = (1/3, 1/3, 1/3)$. The value $\min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta_u(a, A)) = 0.45$ is obtained only for observation $(x, \{x, y, z\})$. Because $\{x, y, z\}\setminus\{x\} = \{y, z\}$ is nonempty, we can select one of the alternatives in $\{y, z\}$, say, y, and move the utility values within the segment $\alpha(0, 1, 0) + (1 - \alpha)u = ((1 - \alpha)/3, (1 + 2\alpha)/3, (1 - \alpha)/3)$. In order to select the appropriate value of α, we consider the observations (a, A) with $a \neq y\in A$ and the observations (y, A). Among the former, the minimal ratio of the data to the Luce probabilities is obtained for $(x, \{x, y, z\})$, with value $0.45/(1 - \alpha)$. In the latter, the minimal ratio is reached at $(y, \{y, z\})$, with value $(0.4(2 + \alpha))/(1 + 2\alpha)$. Equation $0.45/(1 - \alpha) = (0.4(2 + \alpha))/(1 + 2\alpha)$ yields $\bar{\alpha} = 1/4$, which leads to $v = (1/4, 1/2, 1/4)$. The value $\min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta_v(a, A)) = 0.6$ is obtained for pairs $\{(x, \{x, y, z\}), (z, \{x, z\}), (y, \{y, z\})\}$. Notice that the critical observations of δ_v have the cyclical structure described by Proposition 4, that is, $\{x, y, z\} \cup \{x, z\} \cup \{y, z\} = \{x\} \cup \{z\} \cup \{y\}$, and, as a result, the fraction of data explained by the model of Luce cannot be increased further. We have then found the maximal separation $\langle\lambda^*, \delta_L^*, \varepsilon^*\rangle$, with $\delta_L^* = \delta_v$, $\lambda^* = \min_{(a,A)\in\mathcal{O}} (\rho(a, A)/\delta_L^*(a, A)) = 0.6$, and $\varepsilon^* = (\rho - \lambda^*\delta_L^*)/(1 - \lambda^*)$, that is, $\varepsilon^*(x, X) = 0$, $\varepsilon^*(y, X) = 3/4$, $\varepsilon^*(x, \{x, y\}) = 1/8$, $\varepsilon^*(x, \{x, z\}) = 1$, and $\varepsilon^*(y, \{y, z\}) = 0$.

We now illustrate how Proposition 5 works in the example of Table B.1, with the set of single-crossing preferences zP_1yP_1x, yP_2zP_2x, yP_3xP_3z, and xP_4yP_4z, starting with P_1. The maximal fraction of data explained by P_1 is $\lambda_1 = \min_{A\subseteq X} \rho(m_{P_1}(A), A) = \min\{\rho(z, X), \rho(y, \{x, y\}), \rho(z, \{x, z\}), \rho(z, \{y, z\})\} =$

---

32.   In an example that follows, we show that this is not necessarily the general case.

TABLE B.2. P_{DET} and P_{Tremble}

	x	y	z
$\{x, y, z\}$	0.39	0.55	0.06
$\{x, y\}$	0.6	0.4	
$\{x, z\}$	0.95		0.05
$\{y, z\}$		0.95	0.05

$\min\{0.25, 0.75, 0.3, 0.6\} = 0.25$, where, trivially, $\mu_1(P_1) = 1$. We then consider preference P_2, where we have that $\lambda_2 = \min\{\rho(y, X) + \lambda_1, \rho(y, \{x, y\}), \rho(z, \{x, z\}), \rho(y, \{y, z\}) + \lambda_1\} = \min\{0.6 + 0.25, 0.75, 0.3, 0.4 + 0.25\} = 0.3$ with $\mu_2(P_1) = \lambda_1/\lambda_2 = 5/6$ and $\mu_2(P_2) = 1/6$. For preference P_3, we have that $\lambda_3 = \min\{\rho(y, X) + \lambda_1, \rho(y, \{x, y\}), \rho(x, \{x, z\}) + \lambda_2, \rho(y, \{y, z\}) + \lambda_1\} = \min\{0.6 + 0.25, 0.75, 0.7 + 0.3, 0.4 + 0.25\} = 0.65$, with $\mu_3(P_1) = (\lambda_2/\lambda_3)\mu_2(P_1) = 5/13$, $\mu_3(P_2) = (\lambda_2/\lambda_3)\mu_2(P_2) = 1/13$ and $\mu_2(P_3) = 7/13$. Finally, we have that $\lambda_4 = \min\{\rho(x, X) + \lambda_3, \rho(x, \{x, y\}) + \lambda_3, \rho(x, \{x, z\}) + \lambda_2, \rho(y, \{y, z\}) + \lambda_1\} = \min\{0.15 + 0.65, 0.25 + 0.65, 0.7 + 0.3, 0.4 + 0.25\} = 0.65$, and, hence, $\mu_4 = \mu_3$. Thus, we conclude that the maximal fraction of the data that can be explained by SCRUM is 0.65, with maximal SCRUM δ_{μ_4} and residual behavior $\varepsilon(x, X) = 3/7$, $\varepsilon(y, X) = 4/7$, $\varepsilon(x, \{x, y\}) = 5/7$, $\varepsilon(x, \{x, z\}) = 1$, and $\varepsilon(y, \{y, z\}) = 0$, with critical observations (x, X), $(z, \{x, z\})$, and $(y, \{y, z\})$. Note that the example illustrates that the use of a superset of preferences does not necessarily lead to a strict improvement in the goodness of fit.

We now provide an example to illustrate that the deterministic model and the tremble model do not necessarily identify the same preference relations (see Table B.2). Repeating the aforementioned logic, it is easy to see that the optimal preference relation for the deterministic model is $yP_{\text{DET}}xP_{\text{DET}}z$, whereas the one for the tremble model is $xP_{\text{Tremble}}yP_{\text{Tremble}}z$ with a tremble of $30/137$.

Finally, in Section 6, we conjectured that, given the nature of maximal separations, we can expect them to perform better in the overestimation of low observed choice frequencies, whereas maximum likelihood may perform better on average. We then saw this conjecture reflected in the data. Here, we use a simple example involving a particular data-generating process and the tremble model to provide a more formal illustration of the content of our conjecture and the intuition behind it, and leave its further development for future research.

Suppose that the individual has a preference P, and consider a set of binary menus $m_i = \{x_i, y_i\}$, where $x_i P y_i$. The data-generating process entails the maximization of P except for a small menu-dependent error ε_i of choosing alternative y_i.

The log-likelihood of the data with respect to the tremble model is $\sum_i (1 - \varepsilon_i)\log(1 - \gamma/2) + \sum_i \varepsilon_i \log(\gamma/2)$, and its maximization leads to $(1 - (\gamma/2))/(\gamma/2) = (\sum_i (1 - \varepsilon_i))/(\sum_i \varepsilon_i) = (1 - \bar{\varepsilon})/\bar{\varepsilon}$, where $\bar{\varepsilon}$ is the average observed error. That is, maximum log-likelihood averages out the errors observed across different menus, suggesting a tremble of $\gamma_{\text{ML}} = 2\bar{\varepsilon}$ and, consequently, a choice probability for the inferior alternative equal to $\bar{\varepsilon}$. Now consider the maximal separation of the data. For

a given tremble γ, the only potentially critical observations are those in which the mistake is greatest or least, that is, either $\max_i \varepsilon_i$ or $\min_i \varepsilon_i$. In the first case, the superior alternative has been chosen with probability $1 - \max_i \varepsilon_i$, and the estimated tremble model will, by maximal separation, overestimate this probability. Obviously, the superior alternative in any other menu will be less overestimated and cannot be critical. In the second case, likewise, the inferior alternative has been chosen with probability $\min_i \varepsilon_i$, and its maximal separation specification will overestimate it to a greater degree than any other inferior alternative within the remaining menus. In order to find the maximal separation, we need to equalize these two observations, that is, $(1 - (\gamma/2))/(\gamma/2) = (1 - \max_i \varepsilon_i)/\min_i \varepsilon_i$.

For most data-generating processes, for example, any symmetric distribution of error probabilities, the following condition holds: $(1 - \max_i \varepsilon_i)/\min_i \varepsilon_i > (1 - \bar{\varepsilon})/\bar{\varepsilon}$. Whenever this happens, the estimation of maximal separation will provide an estimated tremble $\gamma_{MS} < \gamma_{ML}$, and will, therefore, better accommodate the most extreme observations. The same logic applies to out-of-sample predictions. Consider a new menu with an error probability equal to ε. If $\gamma_{MS} < \gamma_{ML}$, there are three cases of interest: (i) $\varepsilon < \gamma_{MS}/2$, (ii) $\gamma_{MS}/2 < \varepsilon < \gamma_{ML}/2$, and (iii) $\varepsilon > \gamma_{ML}/2$. In the first and third cases, the estimations fail in the same way. That is, they both overestimate the choice probability of the inferior alternative (in case (i)) or the choice probability of the superior alternative (in case (iii)). Clearly, maximal separation does a better job in the first case, where the data are scarce (the relevant alternative being inferior), whereas maximum likelihood does a better job in the latter cases, and also on average.

Appendix C: Empirical Application: Further Considerations

In this section, we report on the application of the maximal separation approach to REU, and the out-of-sample results involving the 3- and 5-option menus.

The REU model proposed by Gul and Pesendorfer (2006) is a key reference in the stochastic treatment of risk preferences. Here we discuss how to use Proposition 1 in order to obtain its maximal separation using our experimental data set.

For the sake of consistency throughout the analysis in this paper, we impose the requirement that all the relevant expected utility preferences be linear orders. Secondly, given that we are working with binary menus, each particular instance of the REU model can be understood as a probability distribution over the set of all preferences satisfying the standard properties of independence and first-order stochastic dominance. Notice that in our setting (i) independence requires that $l_i P l_j$ if and only if $l_{i+4} P l_{j+4}$ for $i, j \in \{2, 3, 4, 5\}$; and (ii) first-order stochastic dominance requires that $l_5 P l_9$. Thus, a REU instance is merely a probability distribution over the set of linear orders satisfying these conditions.

We can then use Proposition 1 to explain how the maximal separation of the data for REU can be obtained. Consider, first, a case of independence, say, $l_4 P l_5$ if and only if $l_8 P l_9$. This leads to the linearity property of REU, where $\delta(l_4, \{l_4, l_5\}) = \delta(l_8, \{l_8, l_9\})$. However, because $\rho(l_4, \{l_4, l_5\}) = 0.49$ and $\rho(l_8, \{l_8, l_9\}) = 0.64$, this is not observed

in the data. Finding the REU instance closest to these data entails finding a value in the range $0.49 < x < 0.64$ such that $0.49/x = (1 - 0.64)/(1 - x)$, which leads to $x = 0.576$. Then, by setting $\delta(l_4, \{l_4, l_5\}) = \delta(l_8, \{l_8, l_9\}) = 0.576$, we obtain a ρ/δ ratio of 0.85, which means that the maximal separation can explain no more than 85% of the data. It can be verified that the other violations of independence are less severe, and, hence, the bound imposed by independence is 0.85. Now consider the implications of stochastic dominance. This requires that, for every instance of REU, it must be that $\delta(l_5, \{l_5, l_9\}) = 1$. However, we observe that $\rho(l_5, \{l_5, l_9\}) = 0.83$, thus yielding a ρ/δ ratio of 0.83. It turns out, therefore, that this ratio determines the goodness-of-fit measure of the REU model in our data set. Clearly, the fraction of the data explained increases with respect to that explained when using CRRA expected utilities, because the latter involve only a subset of expected utilities.

Because REU has no uniqueness in a finite domain, one can find multiple instances of the model for which 83% of the data are explained. We now construct one such instance. Start with the set of lotteries $\{l_3, l_5, l_7, l_8, l_9\}$ and select the following four linear orders over it: (i) $l_8 P_1 l_3 P_1 l_7 P_1 l_5 P_1 l_9$, (ii) $l_5 P_2 l_8 P_2 l_9 P_2 l_7 P_2 l_3$, (iii) $l_3 P_3 l_5 P_3 l_7 P_3 l_9 P_3 l_8$, and (iv) $l_5 P_4 l_9 P_4 l_7 P_4 l_3 P_4 l_8$. Notice that they all place l_5 above l_9, and, hence, any RUM using them will satisfy stochastic dominance. Notice, also, that the independence relationship involving the lotteries l_3, l_5, l_7, and l_9 is always respected. Assign to the four linear orders the probabilities $pq, p(1 - q), (1 - p)q$, and $(1 - p)(1 - q)$, respectively. For each of these four linear orders, consider two linear orders, one with l_1 at the top and the other with l_1 at the bottom, and assign to each of them the conditional probabilities r and $1 - r$, respectively. For each of these eight linear orders, we now place l_4 either at the top or at the bottom, while respecting independence. That is, for the linear orders constructed on the basis of P_1 and P_2, independence requires that l_4 must be above l_3 and l_5, and, hence, we place it at the top. Similarly, for the linear orders constructed on the basis of P_3 and P_4, we place l_4 at the bottom. Finally, for each of these eight linear orders, create one linear order with l_2 at the top followed by l_6, another with l_6 at the top followed by l_2, another with l_2 at the bottom preceded by l_6, and another with l_6 at the bottom preceded by l_2. Notice that this respects independence for any pair associated with l_2 and l_6. Assign to them the conditional probabilities ts, $t(1 - s), (1 - t)s$, and $(1 - t)(1 - s)$. By direct application of Proposition 1, we can find values of these parameters p, q, r, s, and t, which yield the maximal REU separation value 0.83, using the 32 expected utility linear orders described. For instance, $p = 0.565, q = 0.473, r = 0.705, s = 0.2$, and $t = 0.5$. The nature of the residual stochastic choice function follows directly from this construction and Proposition 1.

Our experimental data set involved the choices from 2-, 3- and 5-option menus. In the main text, we have focused on the binary menus, because we have a relatively large number of data points for each binary menu, that is, about 87 choices for each of the 36. In contrast, each participant was confronted with 36 out of the possible 84 menus of 3 lotteries and 36 out of the possible 126 menus of 5 lotteries, all randomly selected without replacement. This gives averages of 37 and 25 observations in the 3- and 5-option menus, respectively, which means markedly fewer data points per menu. In this appendix, we use the data pertaining to the 3- and 5-option menus to perform

TABLE C.1. Summary statistics for the forecasting results for the 3- and 5-option menus.

	Tremble	Luce	SCRUM-CRRA
First quintile	100%	88%	59%
Second quintile	97%	78%	68%
Third quintile	46%	67%	80%
Fourth quintile	0%	34%	82%
Fifth quintile	0%	14%	49%
Average	49%	56%	68%

another out-of-sample exercise. We take the estimated models for maximal separation and maximum likelihood using the binary data reported in Table 3 and follow the methodology adopted in Section 6, evaluating the predictions of these models and techniques using the observations from the nonbinary menus. As in the main text, we focus on those observations in which both maximal separation and maximum likelihood overestimate the observed choice frequencies, and evaluate the probability of the maximal separation prediction being closer to the data than the maximum likelihood prediction. The excessive number of observations makes it unfeasible to report the results for each observation, as in Table 4.

Table C.1 reports some summary statistics. Focusing on those observations for which both maximal separation and maximum likelihood overestimate the observed choice frequencies, and ranking the observations in ascending order of their observed choice frequencies, the table reports, by quintiles and on average, the frequency with which maximal separation is closer than maximum likelihood to the data. We see that, in general, maximal separation is better than maximum likelihood for the low observed choice frequencies. This is particularly true in the case of Tremble and Luce, but also in SCRUM-CRRA when comparing the first quintile against the fifth one. We also see that, on average, maximal separation does a remarkably good job: It is closer to the observed choices than maximum likelihood in 49%, 56%, and 68% of all the overestimated cases. This may have to do with the fact that, in larger menus, the observed choice probabilities are generally smaller, and, thus, better accommodated by maximal separation. However, given the small number of observations for the 3- and 5-option menus, these conclusions should be taken with a grain of salt.

Appendix D: Inconsistency Indices

Starting with Afriat (1973), there is a literature on measuring deviations of actual behavior with respect to the standard, deterministic, rational choice model. Formally, an inconsistency index can be defined as a mapping $I : \text{SCF} \to \mathbb{R}$ describing the inconsistency of a data set $\rho \in \text{SCF}$ with the standard deterministic model, that is, when the reference model is set as $\Delta = \text{DET}$. Most of the existing inconsistency

TABLE D.1. I_{MS} versus I_{HM}.

	x	y	z
$\{x, y, z\}$	0.25	0.3	0.4
$\{x, y\}$	0.8	0.2	
$\{x, z\}$	0.4		0.6
$\{y, z\}$		0.7	0.3

indices are obtained by means of the minimization of a loss function.[33] We can then analyze the inconsistency index emerging from the maximal separation technique. Using the loss function discussed in Section 6, and the insights obtained in Section 4.1, we have

$$I_{\mathrm{MS}} = 1 - \lambda^* = \min_P \max_A \sum_{\substack{a \in A: \\ \delta_P(a,A)=0}} \rho(a, A).$$

It is important to note that the nature of this index is unique in this literature. To illustrate this more clearly, we now compare it with the well-known Houtman and Maks (1985) inconsistency index, which is the closest to I_{MS}. The Houtman–Maks index measures the degree of inconsistency as the minimal number of observations that would have to be removed for the remainder to be consistent with rational choice. The key difference is that the Houtman–Maks index enables different proportions of data to be removed from different menus of alternatives. Hence, using our notation, we can write the Houtman–Maks index as

$$I_{\mathrm{HM}} = \min_P \sum_A \sum_{\substack{a \in A: \\ \delta_P(a,A)=0}} \rho(a, A).$$

These formulations provide a transparent comparison of the two approaches. Both methods remove data minimally until the surviving data are rationalizable. In the case of a maximal separation, because data must be removed at the same rate across all menus, the index focuses on the most problematic one. In the case of Houtman and Maks, different proportions of data can be removed from different menus, which results in an aggregation across menus.

Table D.1 reports an example of a choice function ρ with three alternatives and with data on all the relevant menus of alternatives. Viewed from the I_{MS} perspective, the data show the optimal preference to be $zPxPy$, whereas from the I_{HM} perspective, it is $xP'yP'z$.

---

33. See Apesteguia and Ballester (2015) for a characterization of this class and for a review of the literature.

References

Afriat, S.N. (1973). "On a System of Inequalities in Demand Analysis: An Extension of the Classical Method." *International Economic Review*, 14, 460–472.

Agranov, Marina and Pietro Ortoleva (2017). "Stochastic Choice and Preferences for Randomization." *Journal of Political Economy*, 125(1), 40–68.

Ahn, David S. and Todd Sarver (2013). "Preference for Flexibility and Random Choice." *Econometrica*, 81(1), 341–361.

Apesteguia, Jose and Miguel A. Ballester (2010). "The Computational Complexity of Rationalizing Behavior." *Journal of Mathematical Economics*, 46(3), 356–363.

Apesteguia, Jose and Miguel A. Ballester (2015). "A Measure of Rationality and Welfare." *Journal of Political Economy*, 123(6), 1278–1310.

Apesteguia, Jose and Miguel A. Ballester (2018). "Monotone Stochastic Choice Models: The Case of Risk and Time Preferences." *Journal of Political Economy*, 126(1), 74–106.

Apesteguia, Jose and Miguel A. Ballester and Jay Lu (2017). "Single-Crossing Random Utility Models." *Econometrica*, 85(2), 661–674.

Barseghyan, Levon, Francesca Molinari, Ted O'Donoghue, and Joshua C. Teitelbaum (2016). "Estimating Risk Preferences in the Field." *Journal of Economic Literature*, 56(2), 501–564.

Block, H.D. and Jacob Marschak (1960). "Random Orderings and Stochastic Theories of Response." In *Contributions to Probability and Statistics: Essays in Honor of Harold Hotelling*, edited by I. Olkin et al., Stanford University Press, pp. 97–132.

Böckenholt, Ulf (2006). "Thurstonian-Based Analyses: Past, Present, and Future Utilities." *Psychometrica*, 71(4), 615–629.

Brady, Richard L. and John Rehbeck (2016). "Menu-Dependent Stochastic Feasibility." *Econometrica*, 84(3), 1203–1223.

Caplin, Andrew, Mark Dean, and Daniel Martin (2011). "Search and Satisficing." *American Economic Review*, 101(7), 2899–2922.

Caplin, Andrew and Mark Dean (2015). "Revealed Preference, Rational Inattention, and Costly Information Acquisition." *American Economic Review*, 105(7), 2183–2203.

Crawford, Ian and Krishna Pendakur (2013). "How Many Types Are There?" *Economic Journal*, 123(567), 77–95.

Dean, Mark and Daniel Martin (2016). "Measuring Rationality With the Minimum Cost of Revealed Preference Violations." *Review of Economics and Statistics*, 98(3), 524–534.

Dayton, C. Mitchell (2003). "Applications and Computational Strategies for the Two-Point Mixture Index." *British Journal of Mathematical and Statistical Psychology*, 56, 1–13.

Dickhaut, John, Aldo Rustichini, and Vernon Smith (2009). "A Neuroeconomic Theory of the Decision Process." *Proceedings of the National Academy of Sciences*, 106(52), 22146–22150.

Echenique, Federico, Sangmok Lee, and Matthew Shum (2011). "The Money Pump as a Measure of Revealed Preference Violations." *Journal of Political Economy*, 119(6), 1201–1223.

Echenique, Federico and Kota Saito (2019). "General Luce Model." *Economic Theory*, 68, 811–826.

Famulari, Melissa (1995). "A Household-Based, Nonparametric Test of Demand Theory." *Review of Economics and Statistics*, 77, 372–383.

Fudenberg, Drew, Ryota Iijima, and Tomasz Strzalecki (2015). "Stochastic Choice and Revealed Perturbed Utility." *Econometrica*, 83(6), 2371–2409.

Fudenberg, Drew, Jon Kleinberg, Annie Liang, and Sendhil Mullainathan (2019). "Measuring the Completeness of Theories." Working paper, MIT.

Fudenberg, Drew and Tomasz Strzalecki (2015). "Dynamic Logit With Choice Aversion." *Econometrica*, 83(2), 651–691.

Gul, Faruk and Wolfgang Pesendorfer (2006). "Random Expected Utility." *Econometrica*, 74(1), 121–146.

Gul, Faruk, Paulo Natenzon, and Wolfgang Pesendorfer (2014). "Random Choice as Behavioral Optimization." *Econometrica*, 82(5), 1873–1912.

Halevy, Yoham, Dotan Persitz, and Lanny Zrill (2018). "Parametric Recoverability of Preferences." *Journal of Political Economy*, 126(4), 1558–1593.

Harless, David H. and Colin F. Camerer (1994). "The Predictive Utility of Generalized Expected Utility Theories." *Econometrica*, 62(6), 1251–1289.

Horan, Sean (2019). "Threshold Luce Rules." Working paper, Université de Montréal.

Houtman, Martijn and J.A.H. Maks (1985). "Determining All Maximal Data Subsets Consistent With Revealed Preference." *Kwantitatieve Methoden*, 19, 89–104.

Kalai, Gil, Ariel Rubinstein, and Ran Spiegler (2002). "Rationalizing Choice Functions by Multiple Rationales." *Econometrica*, 70, 2481–2488.

Liang, Annie (2019). "Inference of Preference Heterogeneity From Choice Data." *Journal of Economic Theory*, 179, 275–311

Liu, Jiawei and Bruce G. Lindsay (2009). "Building and Using Semiparametric Tolerance Regions for Parametric Multinomial Models." *Annals of Statistics*, 37(6A), 3644–3659.

Lu, Jay and Kota Saito (2018). "Random Intertemporal Choice." *Journal of Economic Theory*, 177, 780–815.

Luce, Robert D. (1959). *Individual Choice Behavior: A Theoretical Analysis*. Wiley, New York, NY.

Manzini, Paola and Marco Mariotti (2014). "Stochastic Choice and Consideration Sets." *Econometrica*, 82(3), 1153–1176.

Matejka, Filip and Alisdair McKay (2015). "Rational Inattention to Discrete Choices: A New Foundation for the Multinomial Logit Model." *American Economic Review*, 105(1), 272–98.

Mosteller, Frederick and Philip Nogee (1951). "An Experimental Measurement of Utility." *Journal of Political Economy*, 59, 371–404.

Natenzon, Paulo (2019). "Random Choice and Learning." *Journal of Political Economy*, 127(1), 419–457.

Rudas, Tamas (1999). "The Mixture Index of Fit and Minimax Regression." *Metrika*, 50, 163–172.

Rudas, Tamas, Clifford C. Clogg, and Bruce G. Lindsay (1994). "A New Index of Fit Based on Mixture Methods for the Analysis of Contingency Tables." *Journal of the Royal Statistical Society. Series B*, 56(4), 623–639.

Swofford, James L. and Gerald A. Whitney (1987). "Nonparametric Test of Utility Maximization and Weak Separability for Consumption, Leisure and Money." *Review of Economic and Statistics*, 69, 458–464.

Varian, Hal (1990). "Goodness-of-Fit in Optimizing Models." *Journal of Econometrics*, 46, 125–140.

Webb, Ryan (2019). "The (Neural) Dynamics of Stochastic Choice." *Management Science*, 65(1), 230–255.

Supplementary Data

Supplementary data are available at *JEEA* online.